

# CS7150 Deep Learning

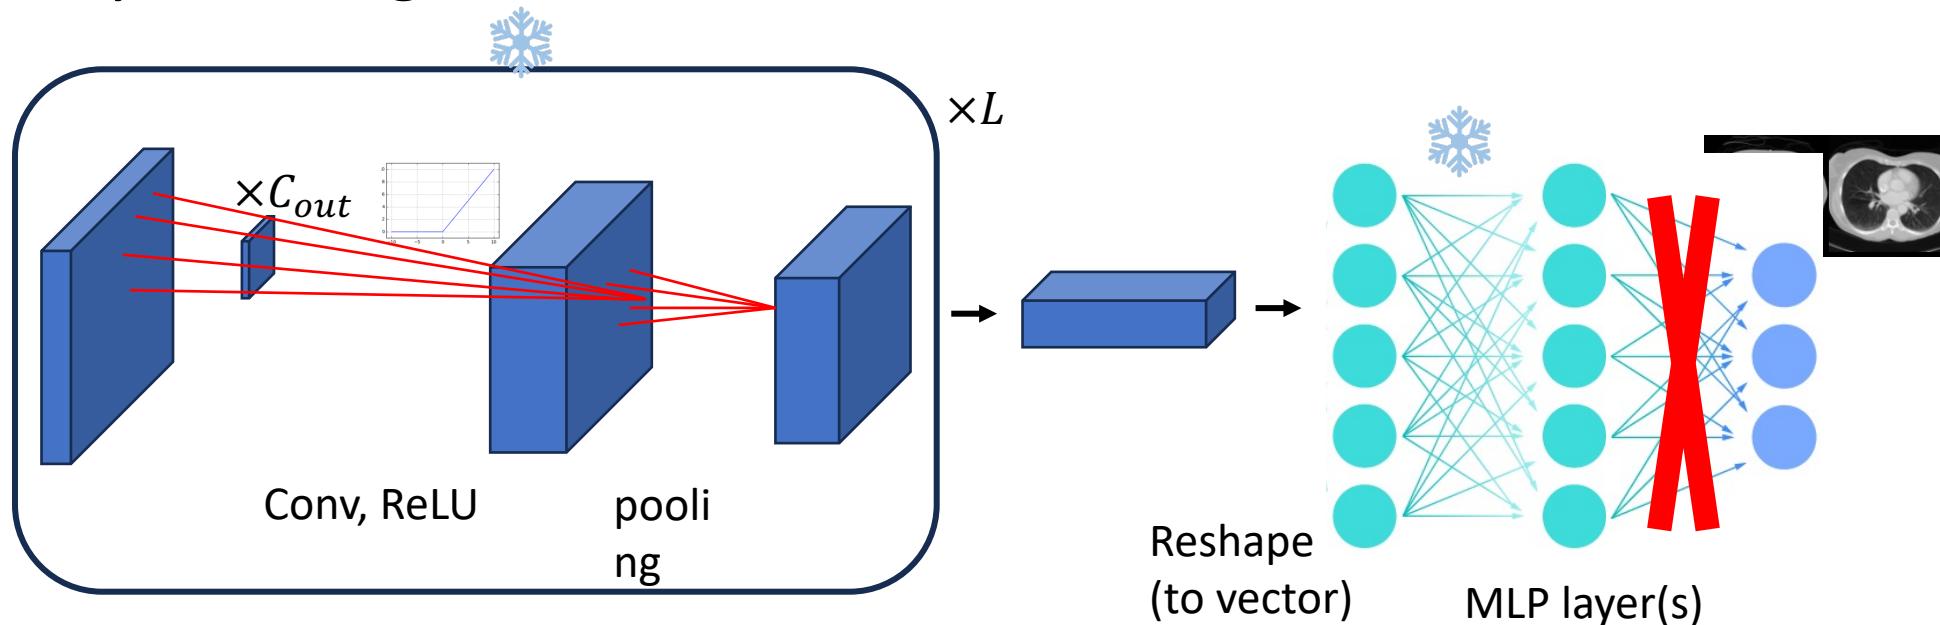
Jiaji Huang

<https://jiaji-huang.github.io>

02/03/2024

# Recap of Last Lecture

- Convolution, correlation, pooling
- Convolutional Neural network
- Transfer learning
- Beyond Image Classification

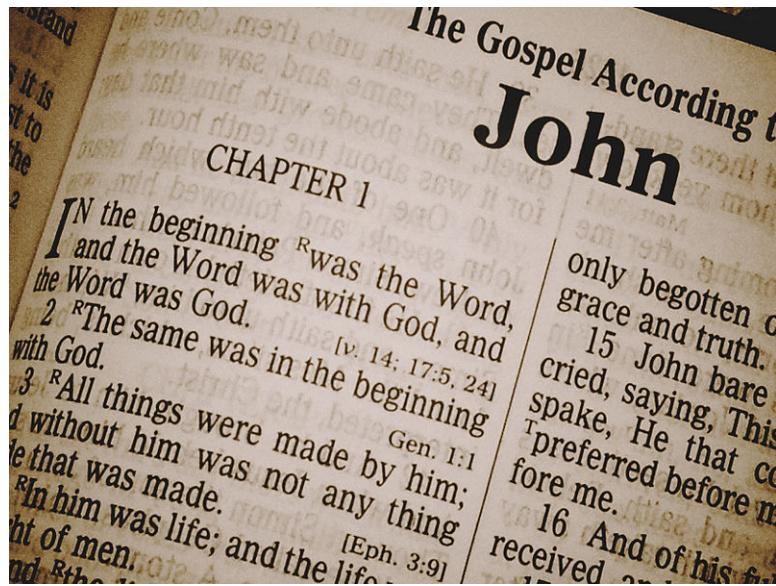


# Agenda

- Learning on Sequences
- Word Embedding
- Recurrent Neural Networks
- Transformer

# Modeling Sequences

- Natural Language



- Source Code

```
<?php language_attributes(); ?>
<?php charset="utf-8"; ?>
<?php bloginfo('charset'); ?>
<?php wp_head(); ?>
<?php wp_title(''); ?>
<?php rel="profile" href="http://gmpg.org/xfn/11"; ?>
<?php rel="pingback" href="<?php bloginfo('pingback_url'); ?>"; ?>
<?php fruitful_get_favicon(); ?>
<?php if (IE_9) ><script src="<?php echo get_template_directory(); ?>/js/ie9.js" type="text/javascript"></script> ?>
<?php wp_head(); ?>
<?php if (is_page()) ><body class="page"> ?>
<?php else() ><body class="not-page"> ?>
<?php if (is_front_page()) ><div id="page-header" class="feed-size"> ?>
<?php $theme_options = fruitful_get_theme_options(); ?>
<?php $logo_pos = $menu_pos = ''; ?>
<?php if (isset($theme_options['logo_pos'])) >$logo_pos = esc_attr($theme_options['logo_pos']); ?>
<?php if (isset($theme_options['menu_pos'])) >$menu_pos = esc_attr($theme_options['menu_pos']); ?>
<?php if (isset($theme_options['menu_class'])) >$menu_class = esc_attr($theme_options['menu_class']); ?>
<?php if (isset($theme_options['menu_sub_class'])) >$menu_sub_class = esc_attr($theme_options['menu_sub_class']); ?>
```

Illustrations from [wiki](#) and this [blogpost](#)

# Modeling Sequences

- Speech Signal

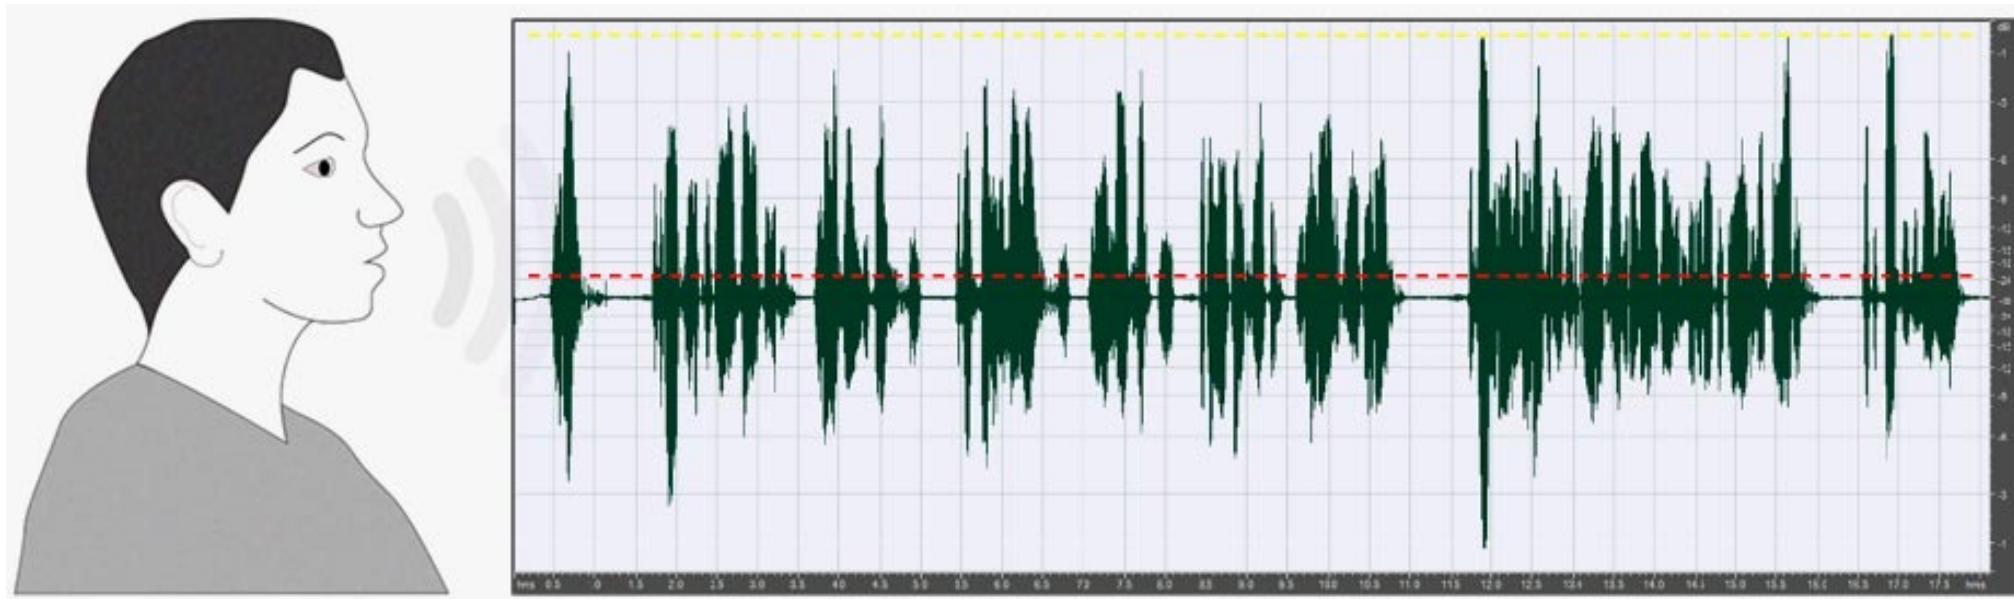


Illustration from this [blogpost](#)

# Modeling Sequences

- DNA

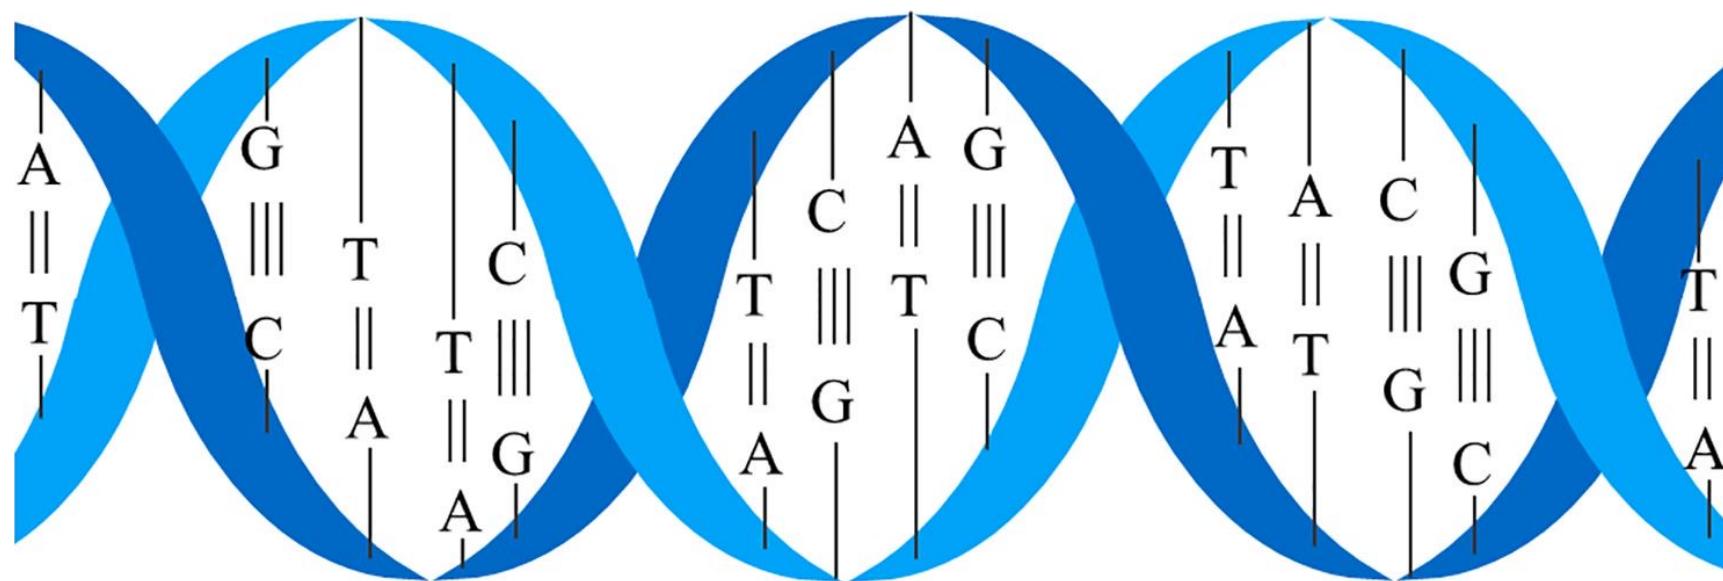


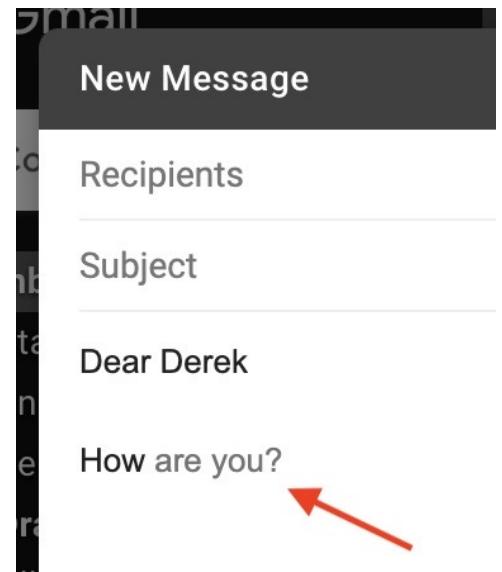
Illustration from [Yang et. al 2020](#)

# Shannon Game

- Guess the next letter (a-z and space) given previous ([C. Shannon, 1950](#))

first line is the original text and the numbers in the second line indicate the guess at which the correct letter was obtained.

- Modern example: auto word suggestion of email clients



# Statistical Language Modeling

- Modeling  $p(w_t | w_1^{t-1})$
- n-gram: assuming  $p(w_t | w_1^{t-1}) = p(w_t | w_{t-(n-1)}^{t-1})$
- unigram  $p(w) = \frac{\#w}{\sum_w \#w}$   
where  $\#w$ : occurrence of word type  $w$  in the corpus
- Bigram  $p(w_t | w_{t-1}) = \frac{\#(w_{t-1}, w_t)}{\#w_{t-1}}$ , enumerate all types of  $(w_{t-1}, w_t)$ 's
- Trigram  $p(w_t | w_{t-2}, w_{t-1}) = \frac{\#(w_{t-2}, w_{t-1}, w_t)}{\#(w_{t-2}, w_{t-1})}$ ,  
enumerate all types of  $(w_{t-2}, w_{t-1}, w_t)$ 's

# “Train” an n-gram Model

- $p(w_t | w_{t-(n-1)}^{t-1}) = \frac{\#(w_{t-(n-1)}, \dots, w_{t-1}, w_t)}{\#(w_{t-(n-1)}, \dots, w_{t-1})}$ , enumerating all types of  $\#(w_{t-(n-1)}, \dots, w_{t-1}, w_t)$
- Gather a corpus
- Calculate unigram  $p(w)$  for all word types  $w$
- For  $k = 2, \dots, n$ :
  - Calculate  $k$ -grams for all types of  $k$ -tuples
  - Store the probabilities as a look-up table
  - Question: what if a  $n$ -tuple type has zero count

# Back-off

- When  $\#(w_{t-(n-1)}, \dots, w_{t-1}, w_t) > \epsilon$

- Discount the usual n-gram

$$p(w_t | w_{t-(n-1)}^{t-1}) = d \cdot \frac{\#(w_{t-(n-1)}, \dots, w_{t-1}, w_t)}{\#(w_{t-(n-1)}, \dots, w_{t-1})}$$

- $p(w_t | w_{t-(n-1)}^{t-1})$  when  $\#(w_{t-(n-1)}, \dots, w_{t-1}, w_t) < \epsilon$

- Back off to  $(n - 1)$  gram

$$p(w_t | w_{t-(n-1)}^{t-1}) = \alpha \cdot p(w_t | w_{t-(n-2)}^{t-1})$$

- $d$  by [Good–Turing estimation](#),  $\alpha$  such that  $p(w_t | w_{t-(n-1)}^{t-1})$  adds to 1

# Test an n-gram Model

- Gather a test corpus (word sequence  $w_1, \dots, w_T$ )
- Calculate

$$p(w_1, \dots, w_T) = \prod_{t=1}^T p(w_t | w_1^{t-1}) = \prod_{t=1}^T p(w_t | w_{\max(1, t-(n-1))}^{t-1})$$

- Very tiny number, use **entropy**:

$$\frac{1}{T} \ln p(w_1, \dots, w_T) = \frac{1}{T} \sum_{t=1}^T \ln p(w_t | w_{\max(1, t-(n-1))}^{t-1})$$

- (Test) **Perplexity**:  $PPL \triangleq \exp \left( -\frac{1}{T} \sum_{t=1}^T \ln p(w_t | w_{\max(1, t-(n-1))}^{t-1}) \right)$

# Entropy and PPL

- If we take base 2
  - The entropy  $\frac{1}{T} \sum_{t=1}^T \log_2 p(w_t | w_{\max(1, t-(n-1))}^{t-1})$  is measured in bits
  - $PPL = 2^{\text{entropy}}$
- $1 \leq PPL \leq \text{all possibilities (vocabulary size)}$
- Back to Shannon game
  - $PPL \leq 27$
  - Therefore Bit per Character (bpc)  $\leq 4.75$
  - As n-gram order increases, we get lower PPL

Column	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	100
Upper	4.03	3.42	3.0	2.6	2.7	2.2	2.8	1.8	1.9	2.1	2.2	2.3	2.1	1.7	2.1	13
Lower	3.19	2.50	2.1	1.7	1.7	1.3	1.8	1.0	1.0	1.0	1.3	1.3	1.2	.9	1.2	.6

Table from ([C. Shannon, 1950](#))

# Packages

- [Kenlm](#), [srilm](#)
- ARPA format

$\log_{10}$  unigram

$\log_{10}$  bigram

```
\data\  
ngram 1=7  
ngram 2=7  
  
\1-grams:  
-1.0000 <unk> -0.2553  
-98.9366 <s> -0.3064  
-1.0000 </s> 0.0000  
-0.6990 wood -0.2553  
-0.6990 cindy -0.2553  
-0.6990 pittsburgh -0.2553  
-0.6990 jean -0.1973  
  
\2-grams:  
-0.2553 <unk> wood  
-0.2553 <s> <unk>  
-0.2553 wood pittsburgh  
-0.2553 cindy jean  
-0.2553 pittsburgh cindy  
-0.5563 jean </s>  
-0.5563 jean wood  
  
\end\
```

Backoff weights ( $\log_{10}$ )

Example from this [blogpost](#)

# Agenda

- Learning on Sequences
- Word Embedding
- Recurrent Neural Networks
- Transformer

# Motivating

- Many more tasks beyond predicting the next word/letter
- E.g., sentiment classification



Cyphis

**[REDACTED] One Star is Too Much for This Product**

Reviewed in the United States on September 7, 2012

I don't know if this is a scam or if mine was broken, but it doesn't work and I am still getting abducted by UFO's on a regular basis.



Amazon Customer

[REDACTED] it works

Reviewed in the United States on March 21, 2018

**Verified Purchase**

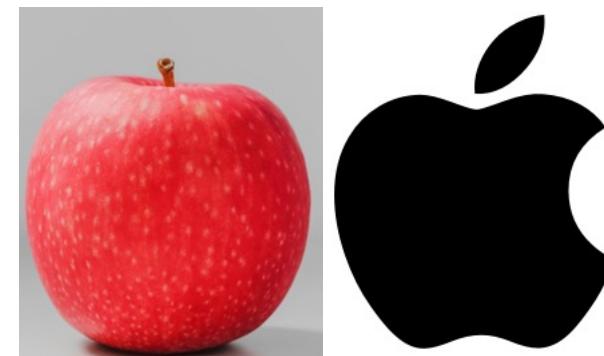
it works great

[UFO detector](#) sold on amazon

# Motivating

- E.g., Named Entity Recognition (NER)

**Apple Surprise Reveal: Biggest-Ever Shake-Up For Your iPhone Here In Weeks**



?

Credit: [Forbes news](#)

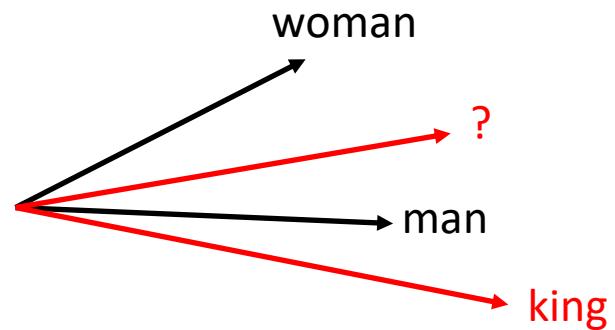
# Word Representations

- Aka embeddings
- One-hot embedding: not very useful
- It's better to encoding some semantics, e.g.,



# Semantics: Word Analogy Task

- Allows us to use linear algebra
- Man to woman is like king to ?



$$\text{man} - \text{woman} = \text{king} - ?$$

$$\Rightarrow ? = \text{king} - \text{man} + \text{woman}$$

$\Rightarrow$  Take ? as the nearest neighbor of RHS

# Word2vec

## Efficient Estimation of Word Representations in Vector Space

by T Mikolov · 2013 · Cited by 37327 — Download a PDF of the paper titled Efficient Estimation of Word Representations in Vector Space, by Tomas Mikolov and 3 other authors.

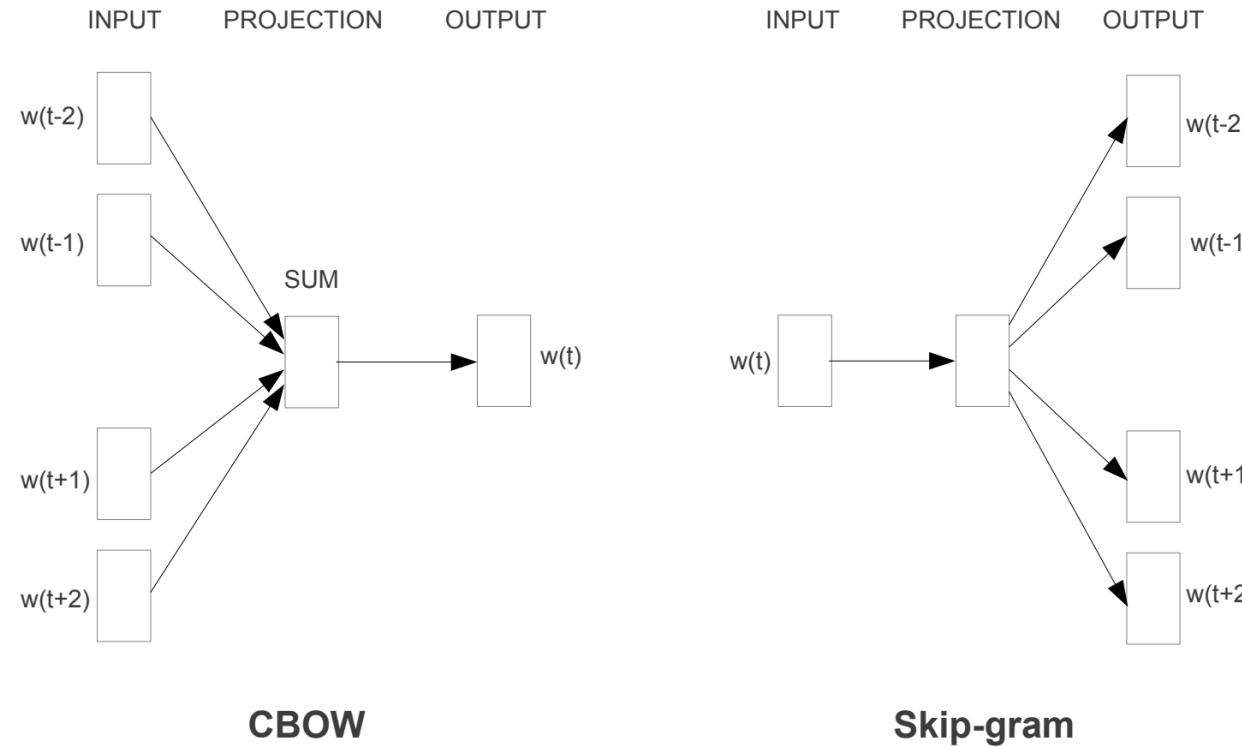


Figure 1: New model architectures. The CBOW architecture predicts the current word based on the context, and the Skip-gram predicts surrounding words given the current word.

# Deep dive into skip-gram

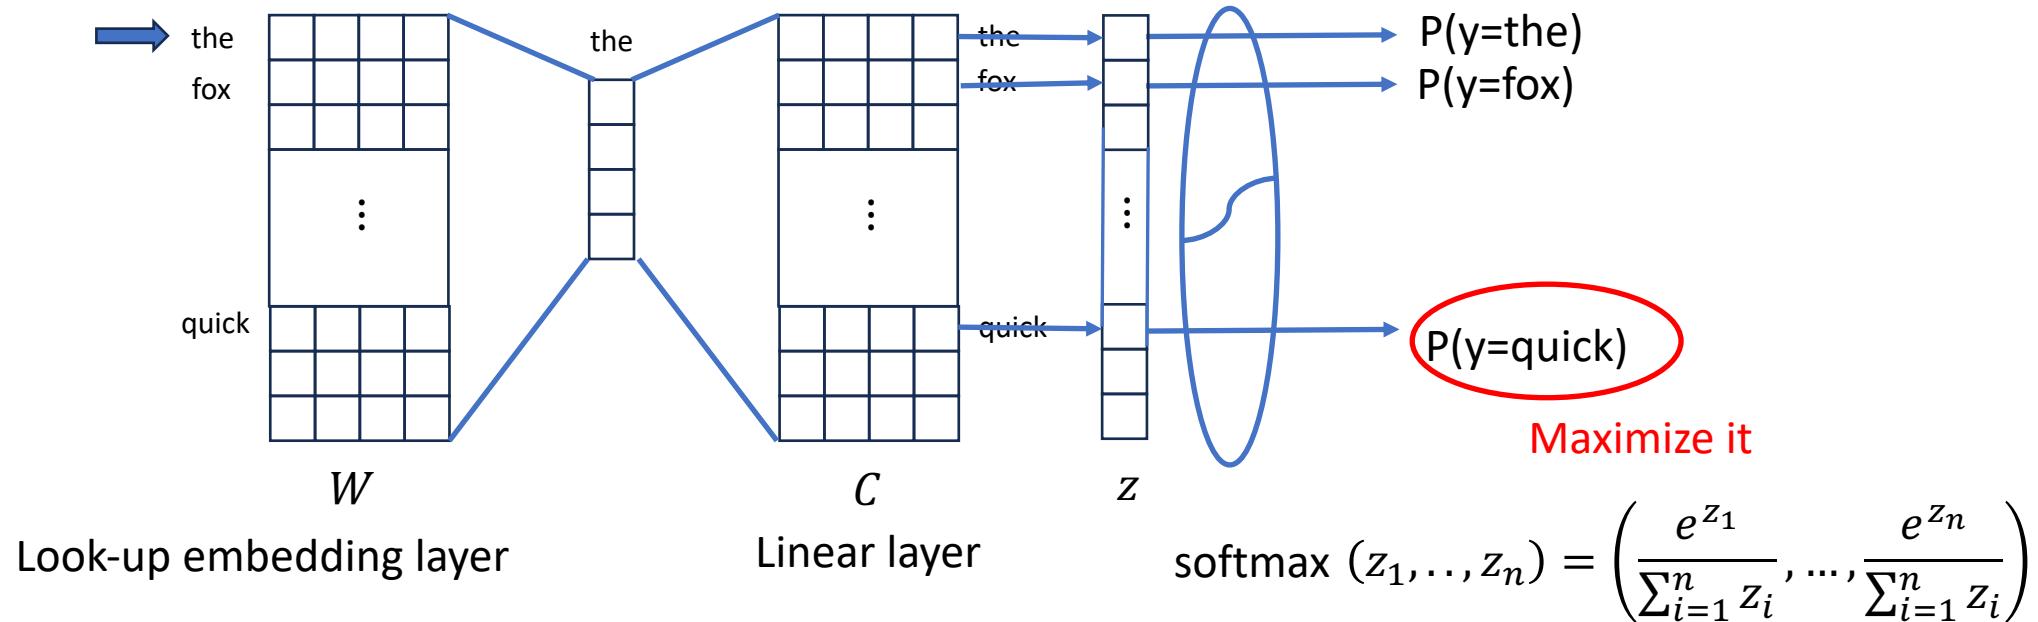
How do we train it? Suppose a window size of 5

Source Text	Training Samples
The quick brown fox jumps over the lazy dog. ➔	(the, quick) (the, brown)
The quick brown fox jumps over the lazy dog. ➔	(quick, the) (quick, brown) (quick, fox)
The quick brown fox jumps over the lazy dog. ➔	(brown, the) (brown, quick) (brown, fox) (brown, jumps)
The quick brown fox jumps over the lazy dog. ➔	(fox, quick) (fox, brown) (fox, jumps) (fox, over)

Figure taken from [blogpost](#)

# Deep dive into skip-gram

- Viewed as 2-layer network
- e.g., receive a training sample ( $w$ =“the”,  $c$ =“quick”)



# Deep dive into skip-gram

- The overall training objective is

$$\min_{W,C} \sum_{w,c} -\log p_{W,C}(c|w)$$

- Typically, we use  $W$  as word representation
- Trained by stochastic gradient descent (SGD)

# Skip-Gram with Negative Sampling (SGNS)



Neural Information Processing Systems

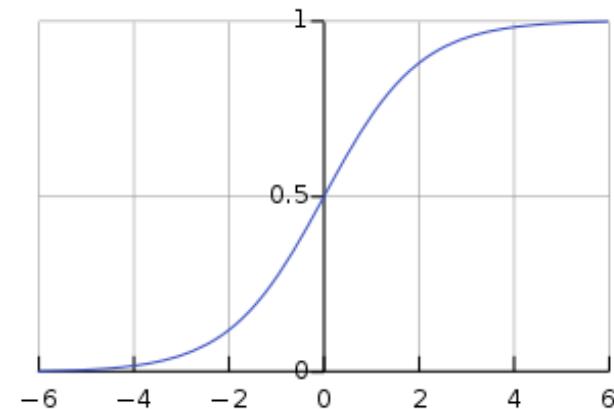
<https://papers.nips.cc/paper/5021-distributed-representations-of-words-and-phrases-in-a-common-embedding-space.pdf>

⋮

[Distributed Representations of Words and Phrases ...](#)

by T Mikolov · 2013 · Cited by 40525 — An inherent limitation of word representations is their indifference to word order and their inability to represent idiomatic phrases. For example, the...

- Vocab can be huge
- huge summation on softmax denominator!
- An efficient variant: contrast with negative samples
- $P(x \text{ and } y \text{ co-occur}) = \sigma(W_x^T C_y)$ ,  
where  $\sigma(s) = \frac{1}{1+e^{-s}}$ , sigmoid function



# Skip-Gram with Negative Sampling (SGNS)

- Denote negative samples as  $y_i$ ,  $i = 1, \dots, k$  from  $q(y)$
- Maximize:

$$\log \sigma(W_x^T C_y) + \sum_{i=1}^k [\log \sigma(-W_x^T C_{y_i})]$$

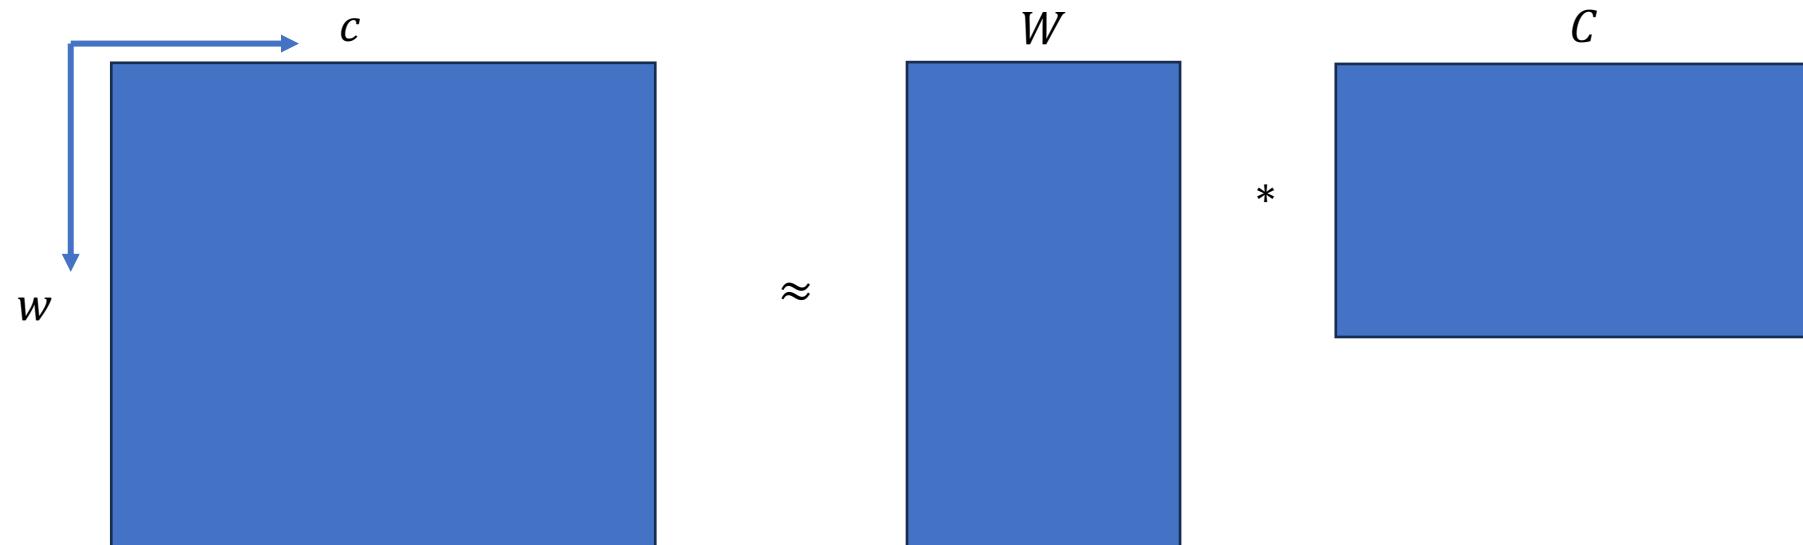
- In practice,  $q(y)$  is uni-gram probability

# Relationship with Matrix Factorization

- ([O. Levy et. al, 2014](#)) shows that SGNS is implicitly factorizing

$$\text{PMI}(w, c) = \log \frac{\#(w, c)}{\#w \cdot \#c}$$

- Glove ([Pennington et. al, 2014](#)): Weighted Factorization of a word-word co-occurrence matrix



# FastText



arXiv

<https://arxiv.org> > cs

:

## Enriching Word Vectors with Subword Information

by P Bojanowski · 2016 · Cited by 11633 — In this [paper](#), we propose a new approach based on the skipgram model, where each word is represented as a bag of character n-grams. A vector ...

- Embedding for a word not seen during training?
- Break word into char n-grams, e.g., n=3  
where → <wh whe her ere re>
- $W["\text{where}"] = W["<\text{wh}>"] + W["\text{whe}"] + W["\text{her}"] + W["\text{ere}"] + W["\text{re}>"]$
- Training: plug above into SGNS
- Testing: for unseen word: sum its char n-gram embeddings

# Evaluation

- Extrinsic evaluation
  - Plug into an ML system and evaluate on a downstream task
  - depends on the ML system
  - But practically very indicative
  - Example: sentiment classification. “I like this movie” v.s. “This is a fair movie”
- Intrinsic evaluation
  - Intermediate, not directly related to downstream task
  - Example: Word analogy

# Why so successful

## Word2Vec

Published online by Cambridge University Press: **16 December 2016**

KENNETH WARD CHURCH

Show author details ▾

Article Figures Metrics



### Abstract

My last column ended with some comments about Kuhn and word2vec. Word2vec has racked up plenty of citations because it satisfies both of Kuhn's conditions for emerging trends: (1) a few initial (promising, if not convincing) successes that motivate early adopters (students) to do more, as well as (2) leaving plenty of room for early adopters to contribute and benefit by doing so. The fact that Google has so much to say on 'How does word2vec work' makes it clear that the definitive answer to that question has yet to be written. It also helps citation counts to [distribute code](#) and data to make it that much easier for the next generation to take advantage of the opportunities (and cite your work in the process).

- native c [implementation](#)
- handy python package: [genism](#), [fastText](#), ...

- A few initial successes
- Plenty of room to improve
- Publish code and data

# Generalize: \*2vec

- Exploit pairwise associations between entities



arXiv

<https://arxiv.org> › cs

⋮

## [node2vec: Scalable Feature Learning for Networks](#)

by A Grover · 2016 · Cited by 11203 — Here we propose **node2vec**, an algorithmic framework for learning continuous feature representations for nodes in networks. In **node2vec**, we learn ...



arXiv

<https://arxiv.org> › cs

⋮

## [\[1902.03545\] Task2Vec: Task Embedding for Meta-Learning](#)

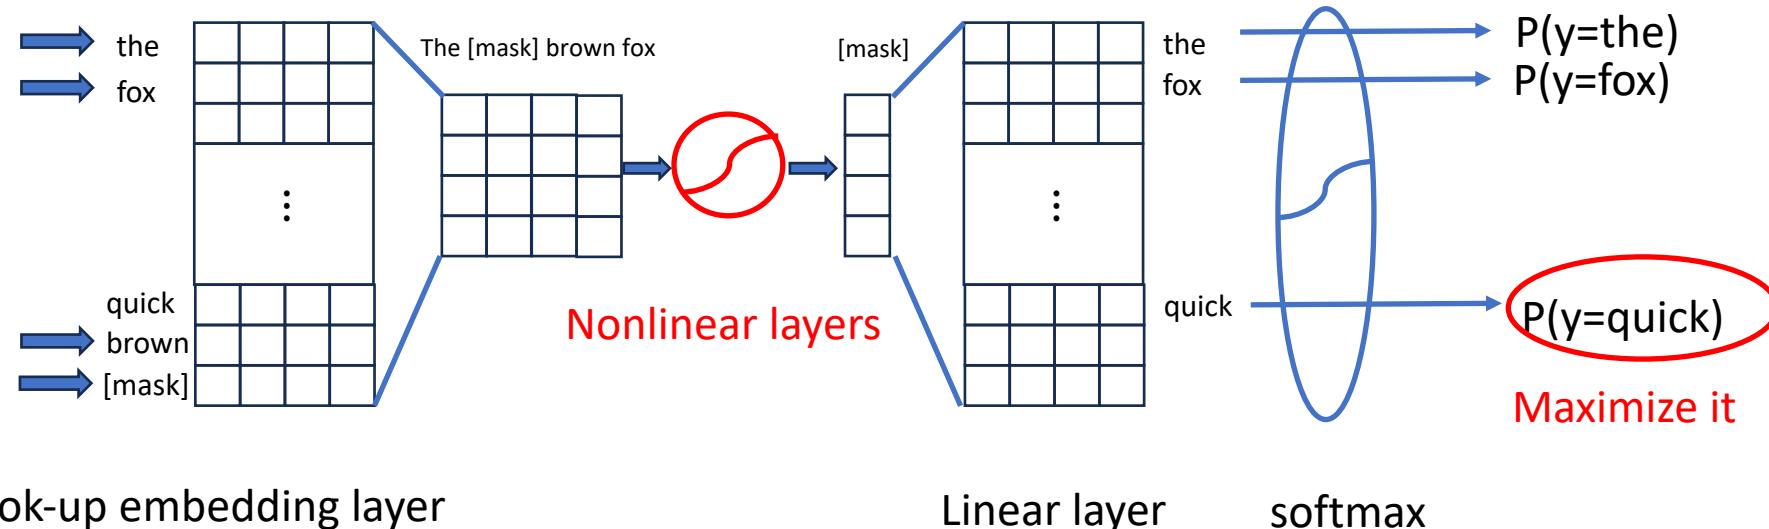
by A Achille · 2019 · Cited by 266 — We present a simple meta-learning framework for learning a metric on embeddings that is capable of predicting which feature extractors will ...

# Agenda

- Learning on Sequences
- Word Embedding
- Recurrent Neural Networks
- Transformer

# Motivation

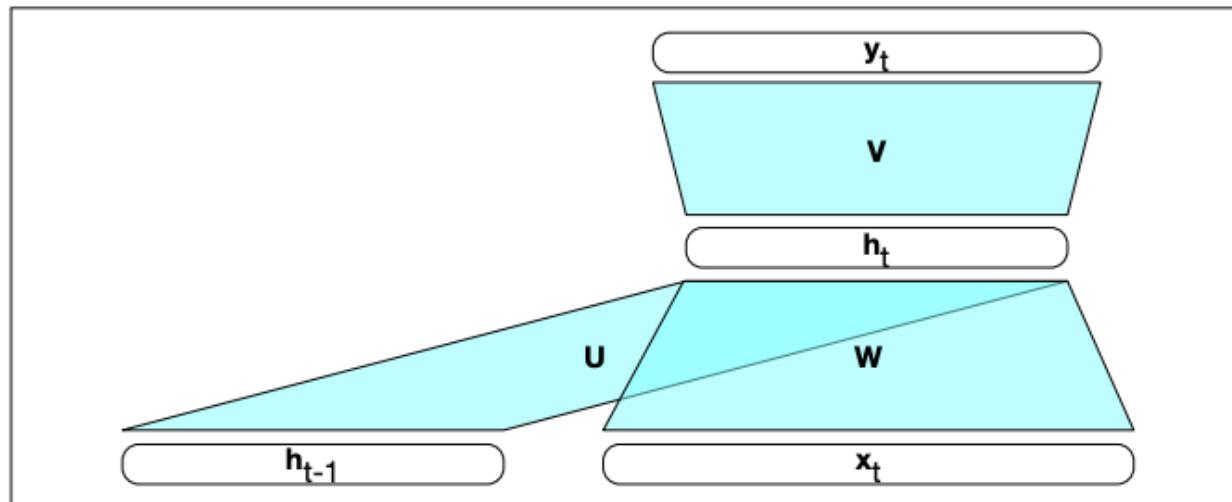
- N-gram and Word2vec both look into a small context window
- Model the entire sequence? E.g.,



- Would MLP work?
- Would conv net work? ([Y. Dauphin et. al, 2017](#))

# Recurrent Neural Network

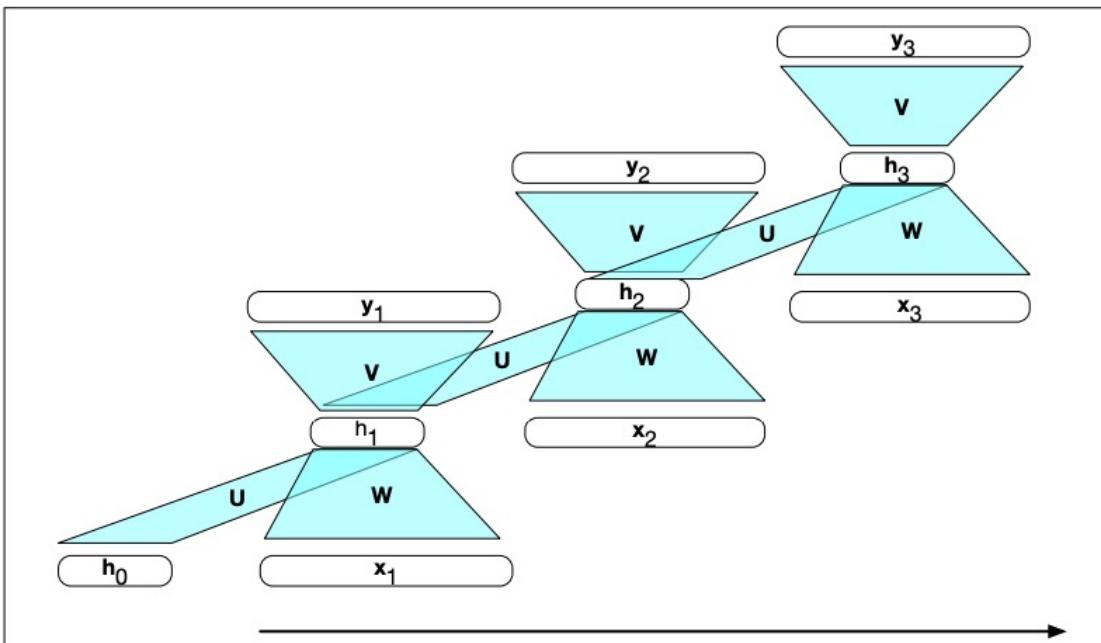
- Hidden state  $h_t$
- Input  $x_t$
- $h_t = g(Uh_{t-1} + Wx_t)$ , e.g.,  $g(\cdot)$  as  $\tanh$
- $y_t = \text{softmax}(Vh_t)$



**Figure 9.2** Simple recurrent neural network illustrated as a feedforward network.

# RNN

- “Unroll” along time axis



**Figure 9.4** A simple recurrent neural network shown unrolled in time. Network layers are recalculated for each time step, while the weights  $U$ ,  $V$  and  $W$  are shared across all time steps.

Illustration from <https://web.stanford.edu/~jurafsky/slp3/9.pdf>

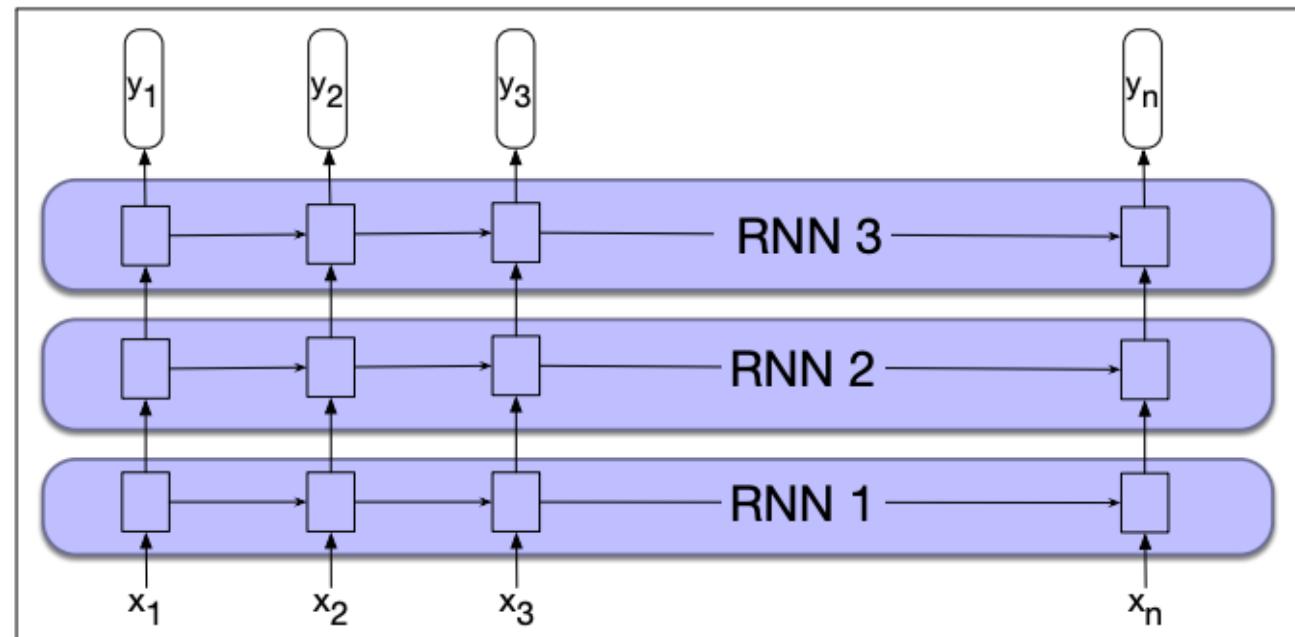
# Stacked RNN layers

$$h_t^1 = g(U^1 h_{t-1}^1 + W^1 x_t)$$

$$h_t^2 = g(U^2 h_{t-1}^2 + W^2 h_t^1)$$

$$h_t^3 = g(U^3 h_{t-1}^3 + W^3 h_t^2)$$

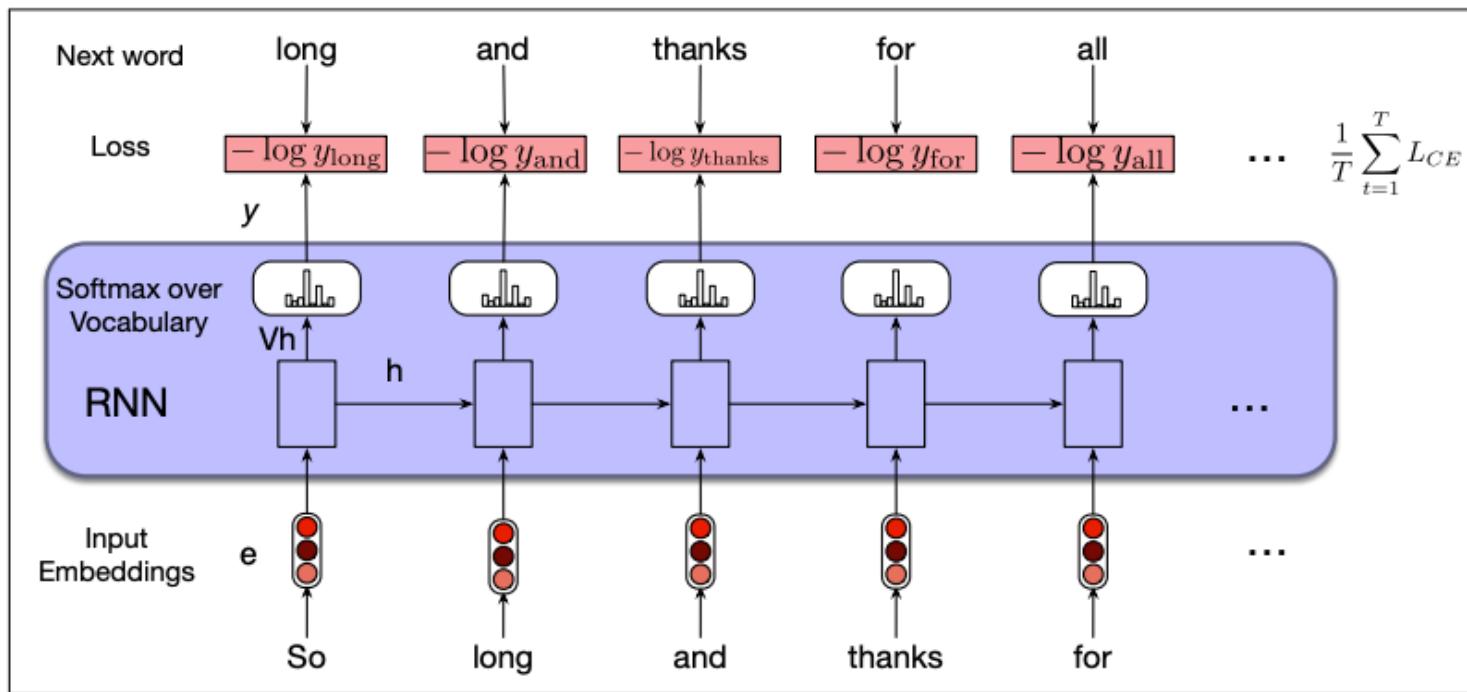
:



**Figure 9.10** Stacked recurrent networks. The output of a lower level serves as the input to higher levels with the output of the last network serving as the final output.

# RNN for language Modeling

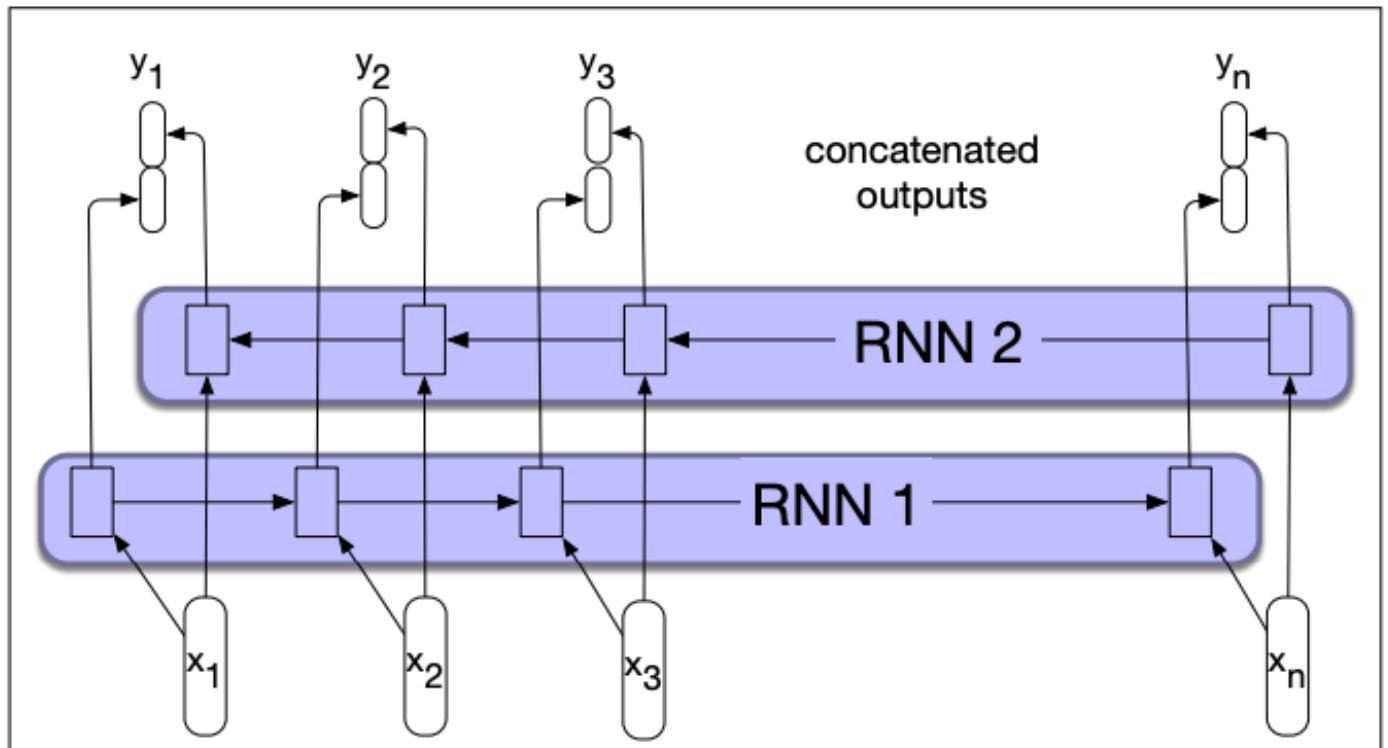
- Predict next word given previous ones



**Figure 9.6** Training RNNs as language models.

# Bi-directional RNN

- Input  $x_1, x_2, \dots, x_T$
- $h_t^f = RNN_f(x_1, x_2, \dots, x_t)$
- $h_t^b = RNN_b(x_T, x_{T-1}, \dots, x_t)$
- $h_t = [h_t^f, h_t^b]$



**Figure 9.11** A bidirectional RNN. Separate models are trained in the forward and backward directions, with the output of each model at each time point concatenated to represent the bidirectional state at that time point.

# Why bi-directional

- Some applications allow us to see the entire sequence, e.g.,
  - Text Classification
  - Speech Recognition
- Bi-directional often beats uni-directional:
  - e.g., a speech recognition example (lower WER is better)

Encoder Architecture	WSJ		HKUST		LibriSpeech			
	dev	test	dev	test	dev	dev	test	test
					clean	other	clean	other
“Google”-BLSTM	7.9	4.7	29.9	28.9	4.7	14.1	4.9	15.2
“Google”-LSTM	9.9	6.5	35.5	33.8	6.3	18.2	6.5	19.4

ASR results from <https://www.jonathanleroux.org/pdf/Moritz2019Interspeech09.pdf>

# LSTM, but why?

- RNNs can be hard to train
- Gradient explosion and vanishing problem



arXiv

<https://arxiv.org> > cs

:

## On the difficulty of training Recurrent Neural Networks

by R Pascanu · 2012 · Cited by 6493 — We propose a gradient norm clipping strategy to deal with exploding gradients and a soft constraint for the vanishing gradients problem. We ...

- In *modern words*, attention span is very short!

# LSTM

- Introduces memory  $c_t$
- Forget gate

$$f_t = \sigma(U_f h_{t-1} + W_f x_t)$$

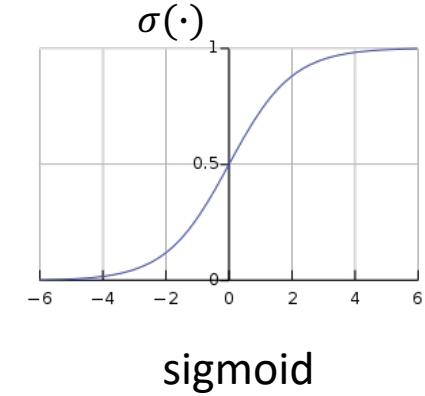
- Add gate

$$\begin{aligned} i_t &= \sigma(U_i h_{t-1} + W_i x_t) \\ g_t &= \tanh(U_g h_{t-1} + W_g x_t) \end{aligned}$$

- Combined
- Output gate

$$c_t = g_t \odot i_t + c_{t-1} \odot f_t$$

$$\begin{aligned} o_t &= \sigma(U_o h_{t-1} + W_o x_t) \\ h_t &= o_t \odot \tanh(c_t) \end{aligned}$$



# GRU

- A Variant of LSTM, no memory cells
- fewer parameters

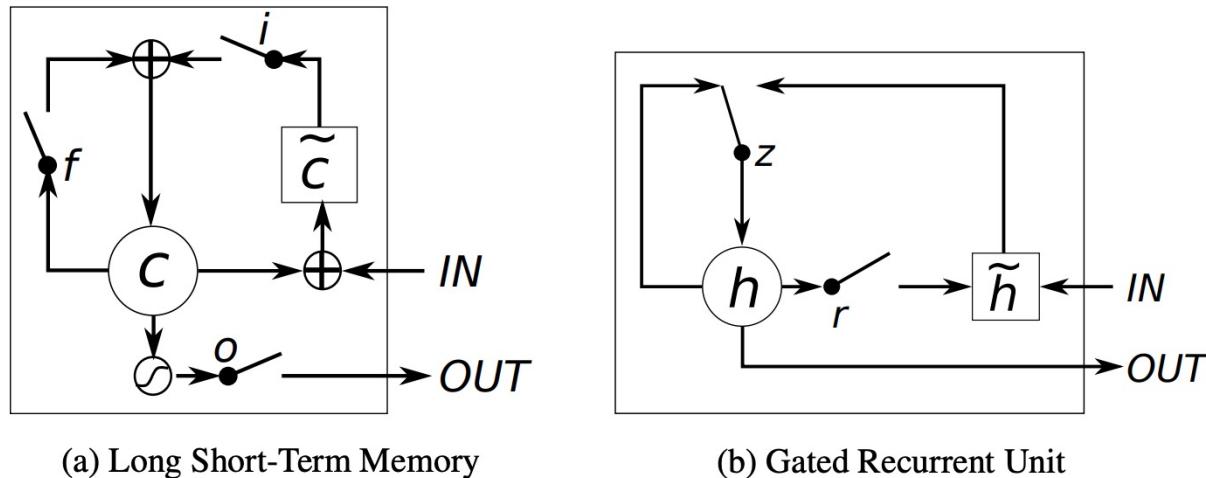


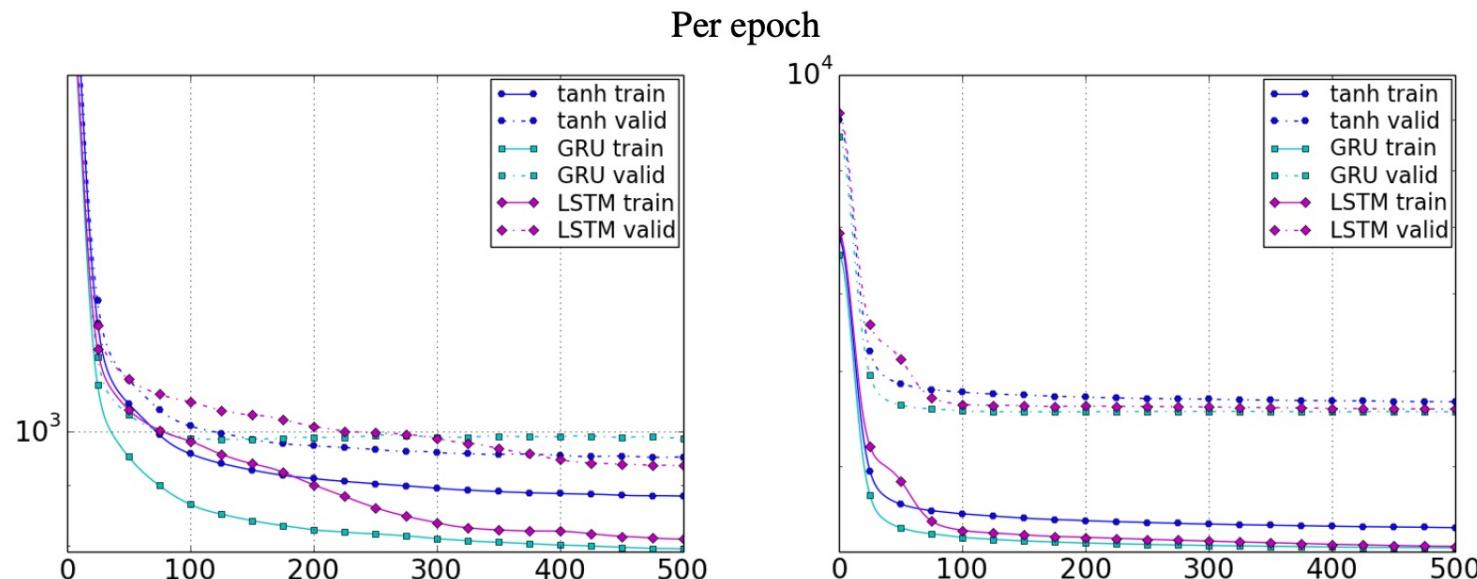
Figure 1: Illustration of (a) LSTM and (b) gated recurrent units. (a)  $i$ ,  $f$  and  $o$  are the input, forget and output gates, respectively.  $c$  and  $\tilde{c}$  denote the memory cell and the new memory cell content. (b)  $r$  and  $z$  are the reset and update gates, and  $h$  and  $\tilde{h}$  are the activation and the candidate activation.

# Comparison

## Empirical Evaluation of Gated Recurrent Neural Networks ...

by J Chung · 2014 · Cited by 13908 — In this paper we compare different types of recurrent units in recurrent neural networks (RNNs). Especially, we focus on more sophisticated ...

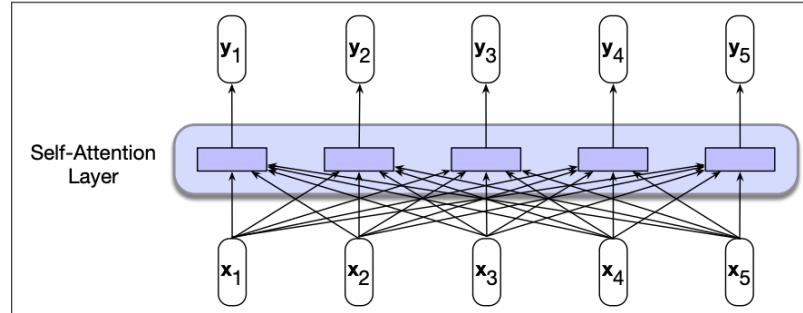
These results clearly indicate the advantages of the gating units over the more traditional recurrent units. Convergence is often faster, and the final solutions tend to be better. However, our results are not conclusive in comparing the LSTM and the GRU, which suggests that the choice of the type of gated recurrent unit may depend heavily on the dataset and corresponding task.



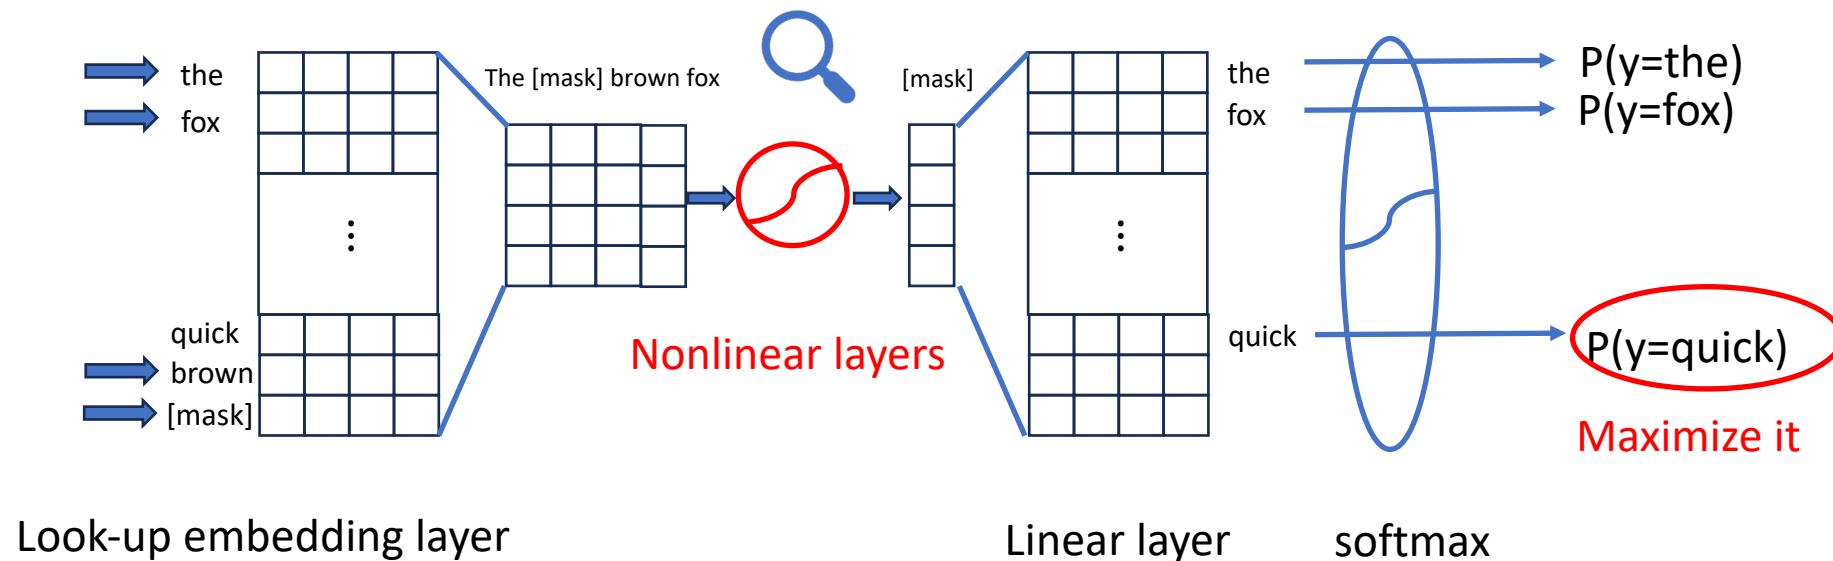
# Agenda

- Learning on Sequences
- Word Embedding
- Recurrent Neural Networks
- Transformer

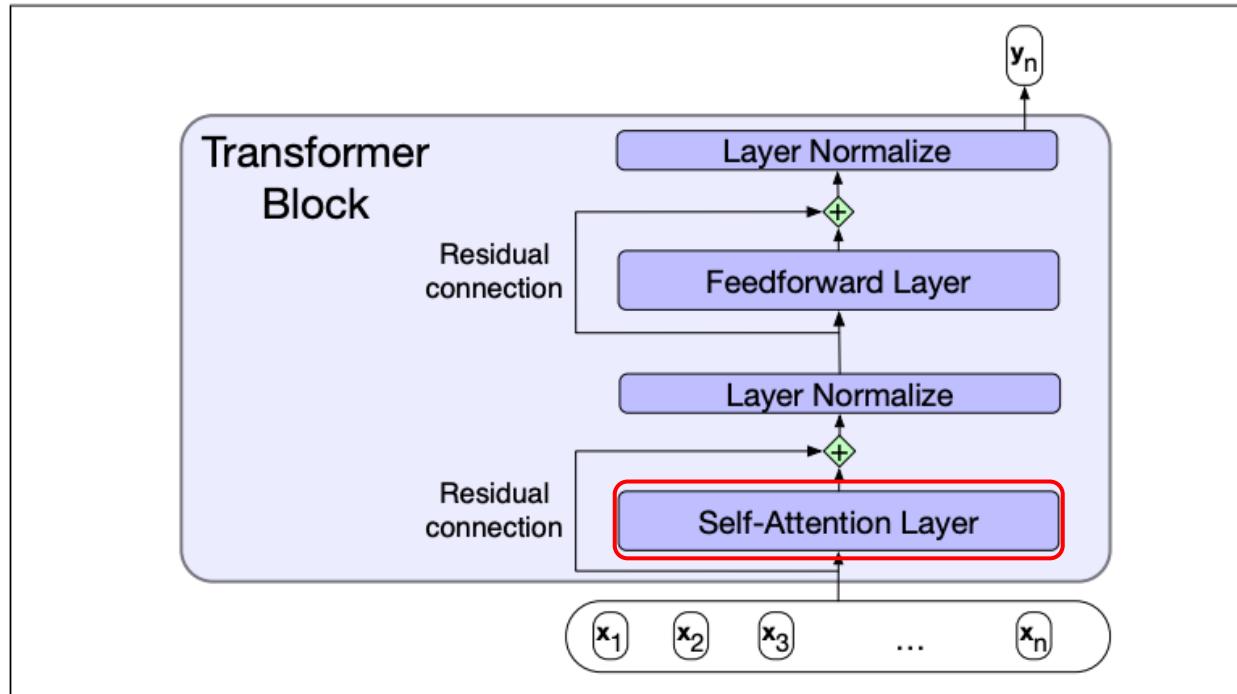
# Motivating



**Figure 11.2** Information flow in a bidirectional self-attention model. In processing each element of the sequence, the model attends to all inputs, both before and after the current one.



# A Transformer Block



**Figure 10.4** A transformer block showing all the layers.

Illustration from <https://web.stanford.edu/~jurafsky/slp3/10.pdf>

# Math of Self-attention

- Map hidden states to keys, queries and values ( $d$  dimensional)

$$q_i = W^q x_i, k_i = W^k x_i, v_i = W^v x_i$$

- Raw score

$$s_{i,j} = q_i \cdot k_j / \sqrt{d}$$

- Attention weights

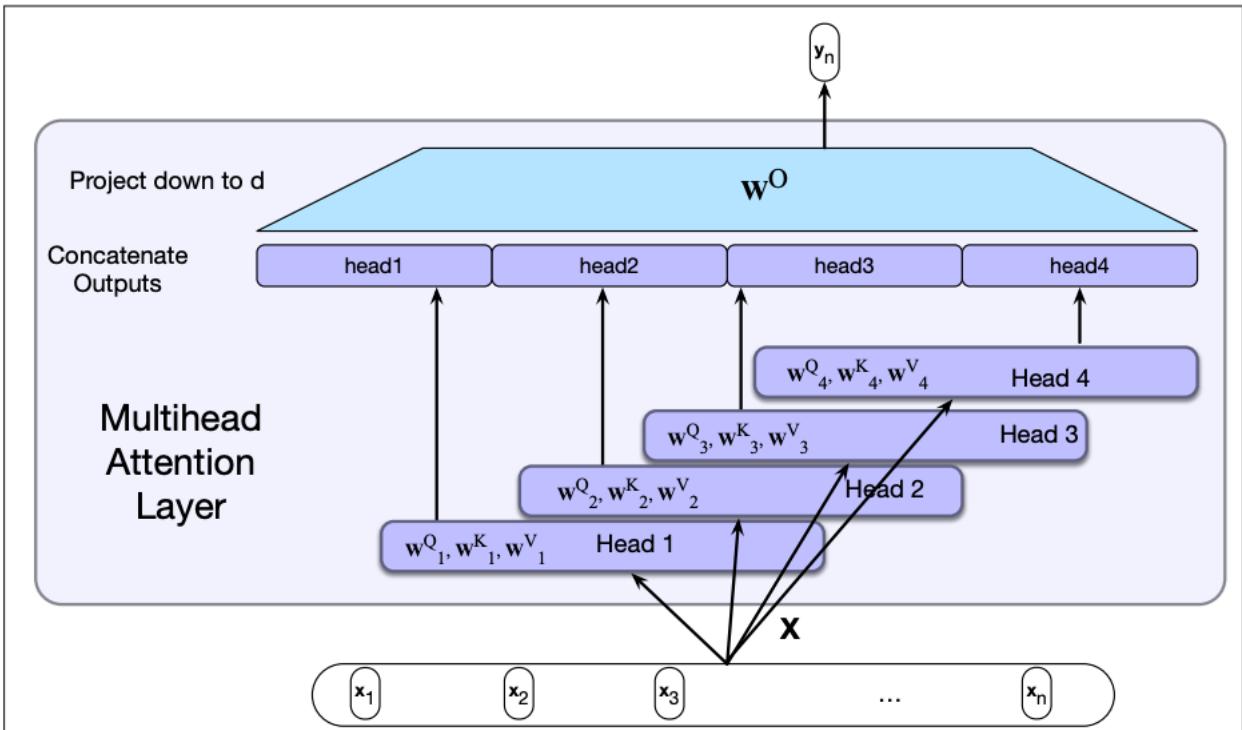
$$\alpha_{i,j} = \frac{e^{s_{i,j}}}{\sum_l e^{s_{i,l}}}$$

- Output Hidden states

$$y_i = \sum_j \alpha_{i,j} v_j$$

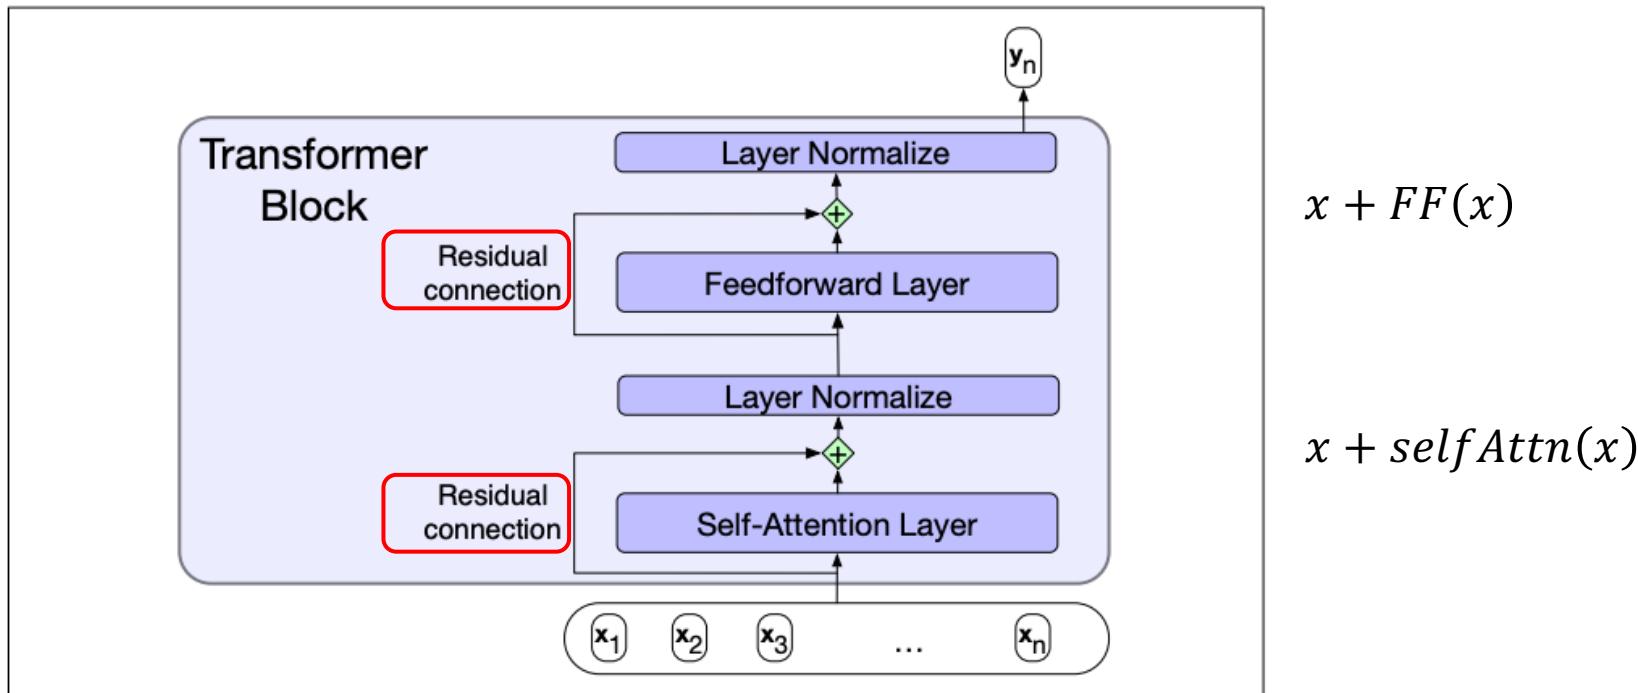
# Multi-head Attention

- Input  $x \in \mathbb{R}^d$
- Multiple sets of  $W^q$ ,  $W^k$  and  $W^v$ 
  - Each project to  $d/\#\text{heads}$
  - Typically,  
 $\#\text{heads}=8, 12, 16, 24, \dots$
- concatenate output from each head
- Project with  $W^o$ , usually same  $d$



**Figure 10.5** Multihead self-attention: Each of the multihead self-attention layers is provided with its own set of key, query and value weight matrices. The outputs from each of the layers are concatenated and then projected down to  $d$ , thus producing an output of the same size as the input so layers can be stacked.

# Residual Connection (refer to last lecture)

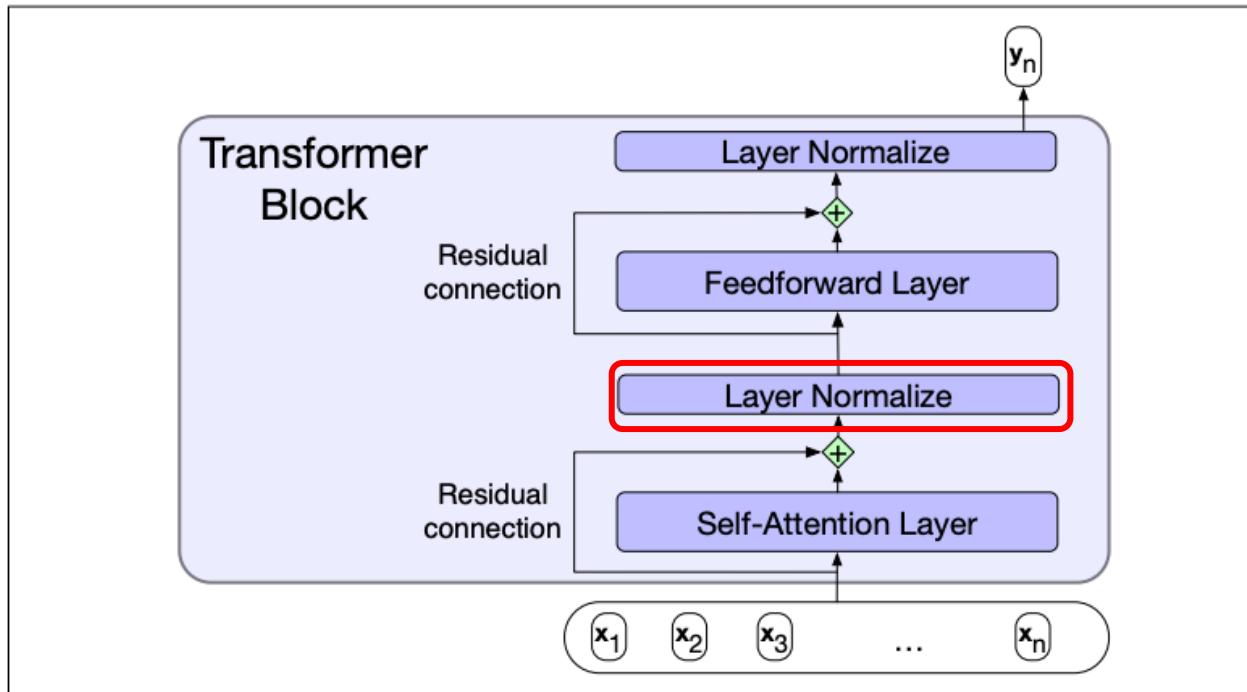


**Figure 10.4** A transformer block showing all the layers.

Illustration from <https://web.stanford.edu/~jurafsky/slp3/10.pdf>

## [1607.06450] Layer Normalization

by JL Ba · 2016 · Cited by 9138 — Layer normalization is very effective at stabilizing the hidden state dynamics in recurrent networks. Empirically, we show that layer ...

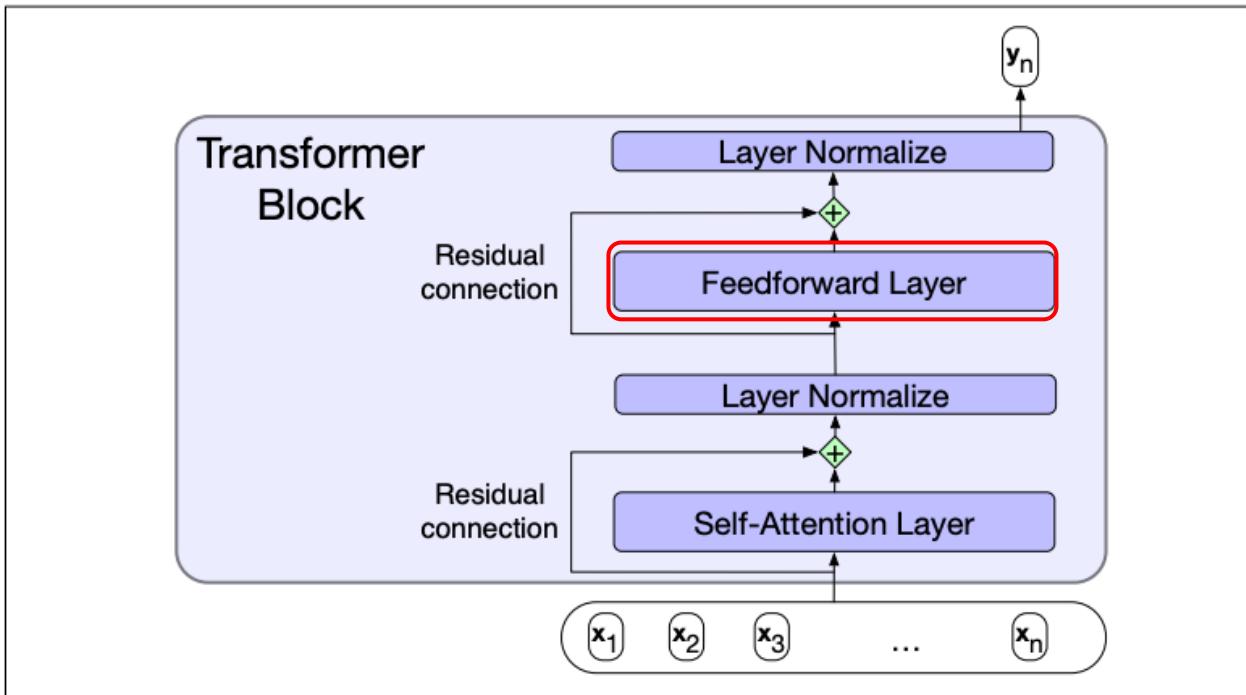


**Figure 10.4** A transformer block showing all the layers.

$$\mu = \frac{1}{d_h} \sum_{i=1}^{d_h} x_i$$
$$\sigma = \sqrt{\frac{1}{d_h} \sum_{i=1}^{d_h} (x_i - \mu)^2}$$
$$\hat{\mathbf{x}} = \frac{(\mathbf{x} - \mu)}{\sigma}$$

$$\text{LayerNorm} = \gamma \hat{\mathbf{x}} + \beta$$

# Feedforward Layer



**Figure 10.4** A transformer block showing all the layers.

$$y = W_2 \sigma(W_1 x)$$

intermediate dimension: Usually  $4 \times$  input dimension

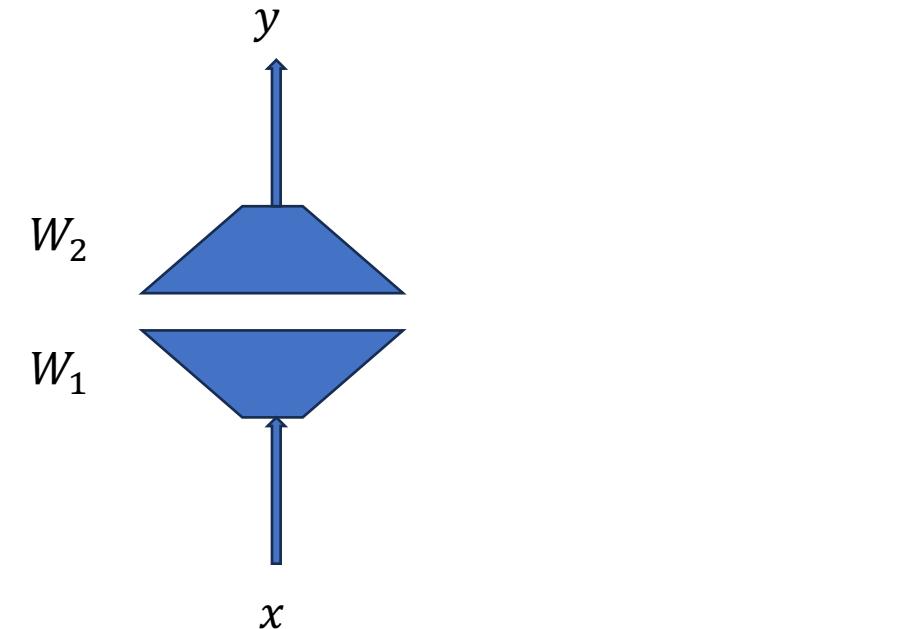
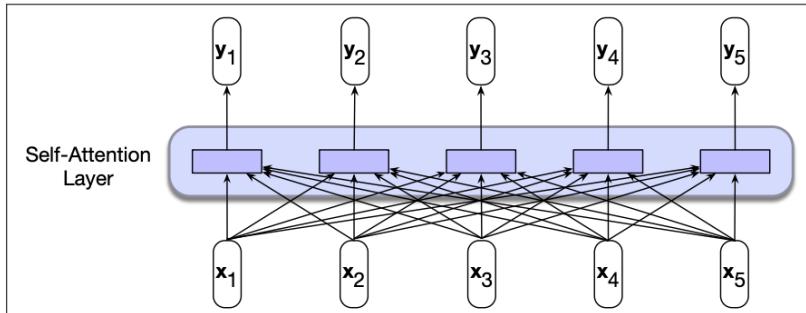
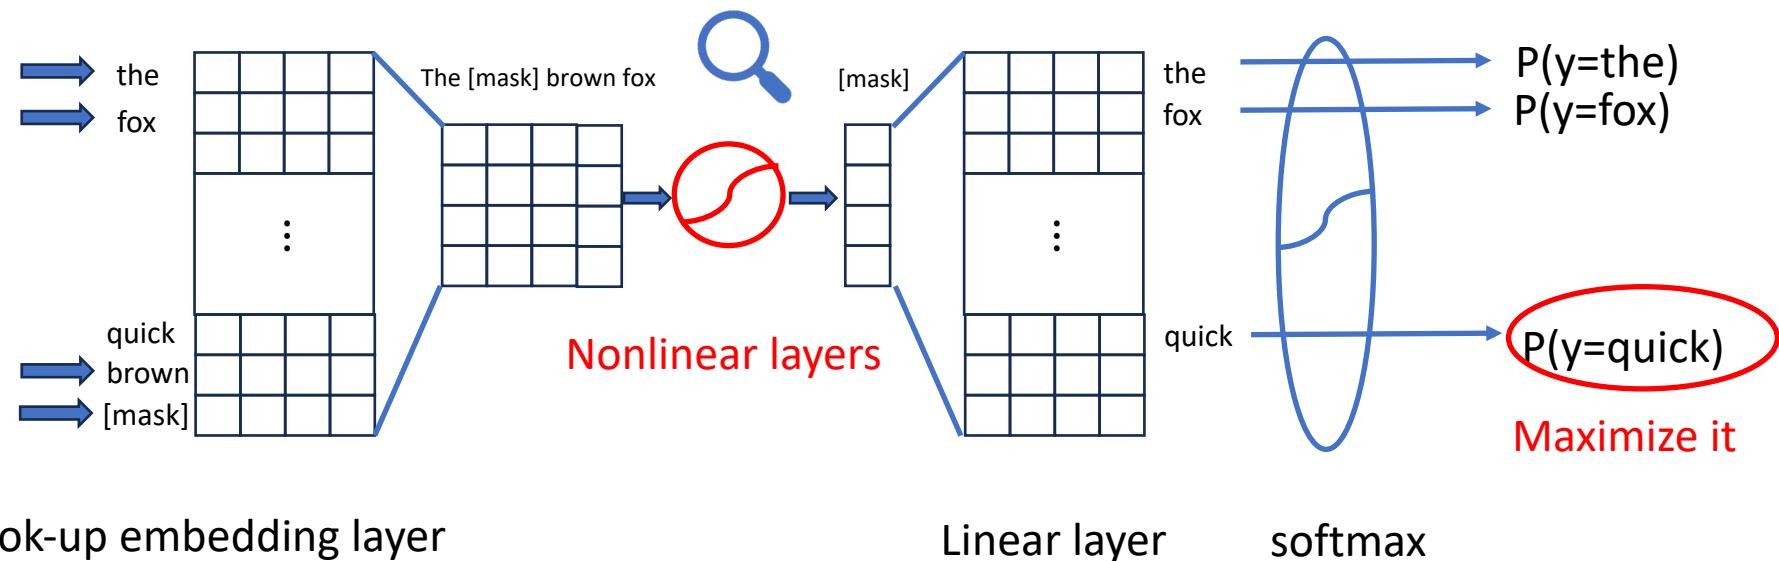


Illustration from <https://web.stanford.edu/~jurafsky/slp3/10.pdf>

# Transformer Encoder for masked LM



**Figure 11.2** Information flow in a bidirectional self-attention model. In processing each element of the sequence, the model attends to all inputs, both before and after the current one.



# (Causal) Attention Mask

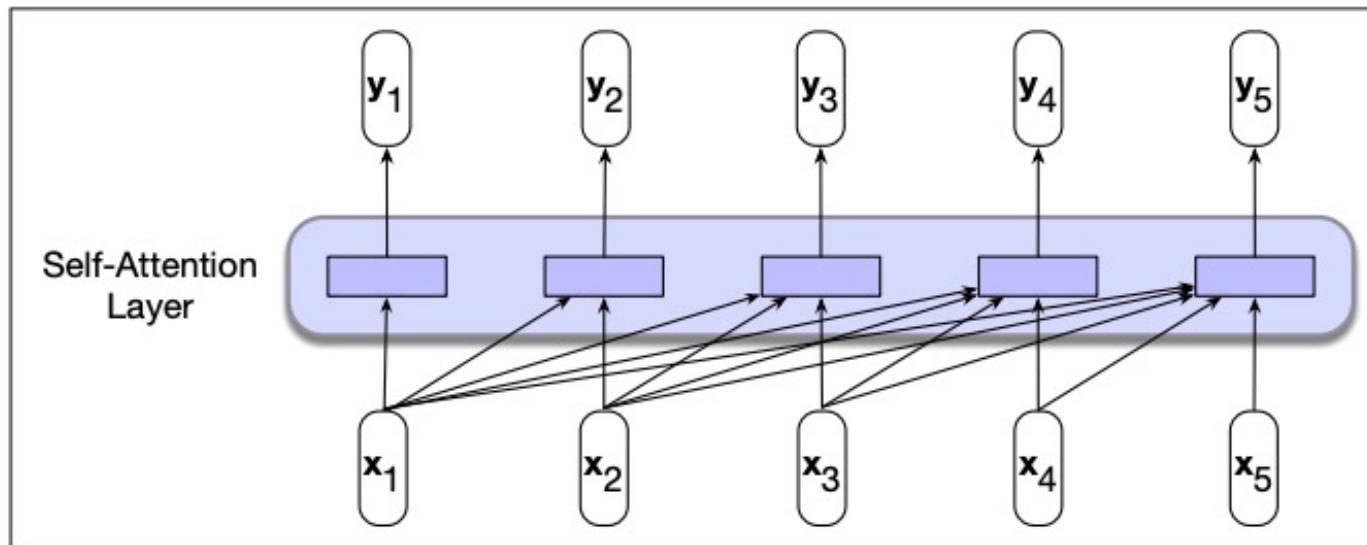
- $\alpha_{i,j} = \frac{e^{s_{i,j}}}{\sum_l e^{s_{i,l}}}, y_i = \sum_j \alpha_{i,j} v_j$
- The range of  $j$  (or  $l$ ) for causal language model:  $j \leq i$

N	q1·k1	-∞	-∞	-∞	-∞
	q2·k1	q2·k2	-∞	-∞	-∞
	q3·k1	q3·k2	q3·k3	-∞	-∞
	q4·k1	q4·k2	q4·k3	q4·k4	-∞
	q5·k1	q5·k2	q5·k3	q5·k4	q5·k5

**Figure 10.3** The  $N \times N$   $\mathbf{Q}\mathbf{K}^T$  matrix showing the  $q_i \cdot k_j$  values, with the upper-triangle portion of the comparisons matrix zeroed out (set to  $-\infty$ , which the softmax will turn to zero).

# Transformer Decoder for (Causal) LM

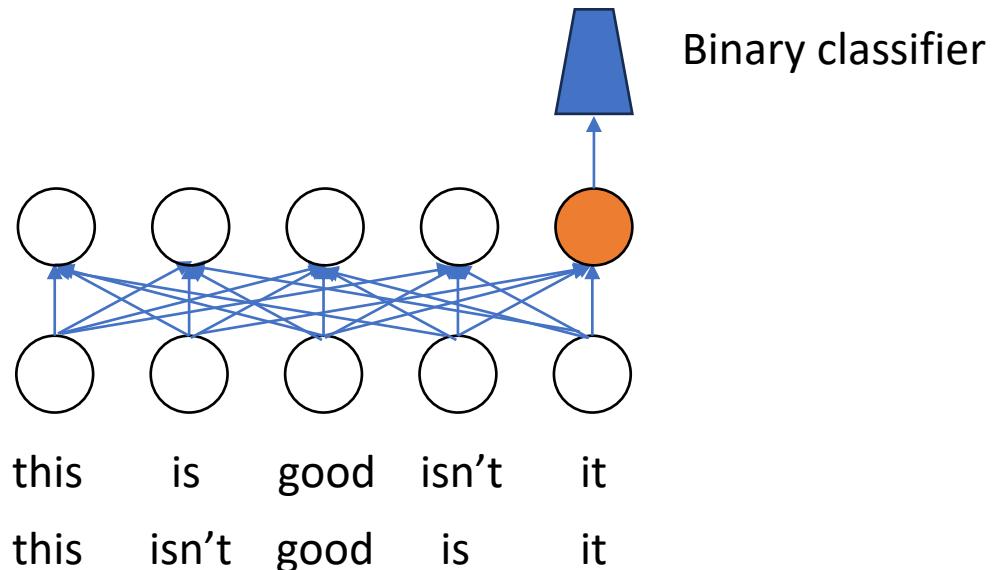
- With Attention mask, we can predict next word given previous ones



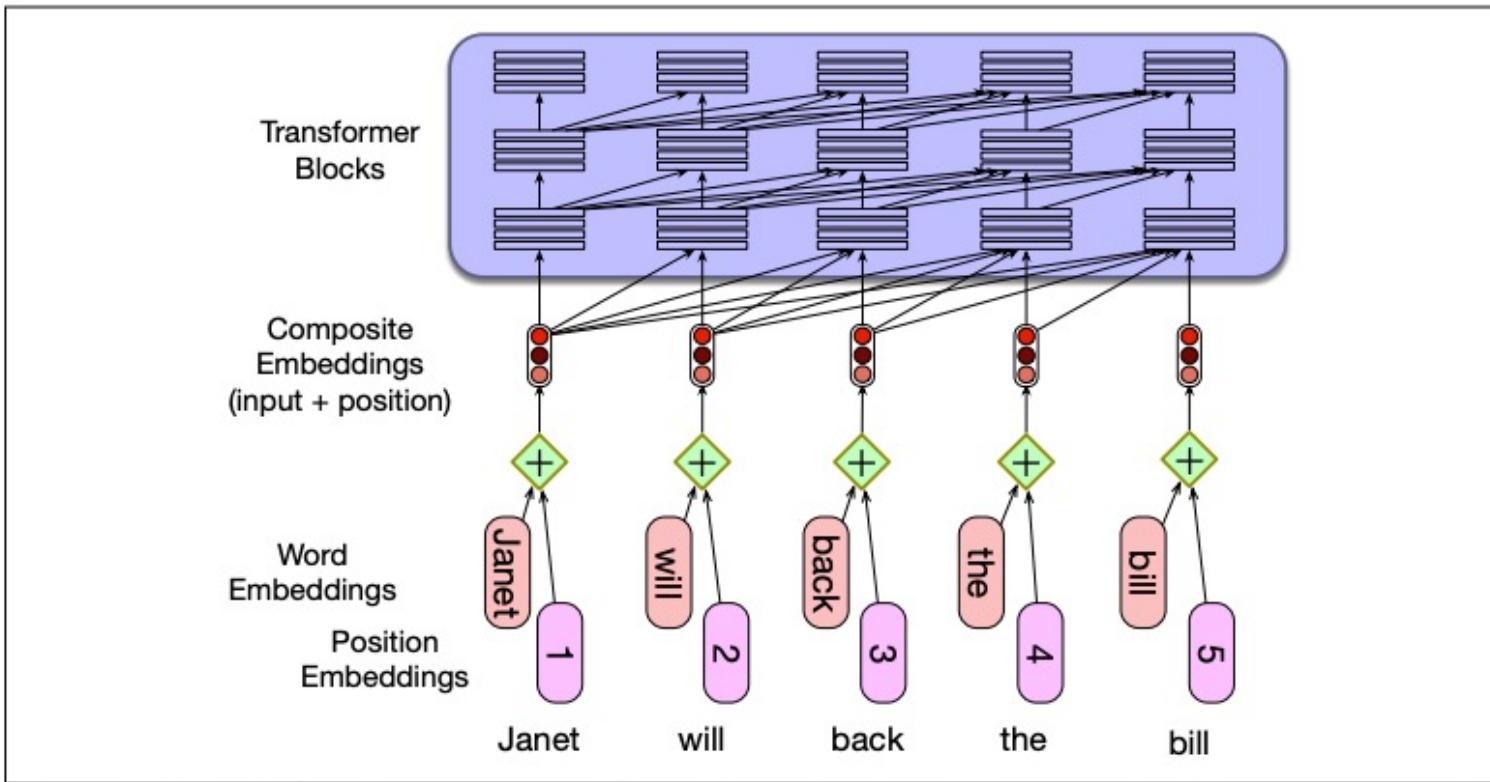
**Figure 10.1** Information flow in a causal (or masked) self-attention model. In processing each element of the sequence, the model attends to all the inputs up to, and including, the current one. Unlike RNNs, the computations at each time step are independent of all the other steps and therefore can be performed in parallel.

# Did we miss anything?

- Consider sentiment classification  
“this is good, isn’t it?”  
vs  
“this isn’t good, is it?”
- Same hidden state for classification
- We’re agnostic to word orders!



# Positional Encoding



**Figure 10.6** A simple way to model position: simply adding an embedding representation of the absolute position to the input word embedding to produce a new embedding of the same dimensionality.

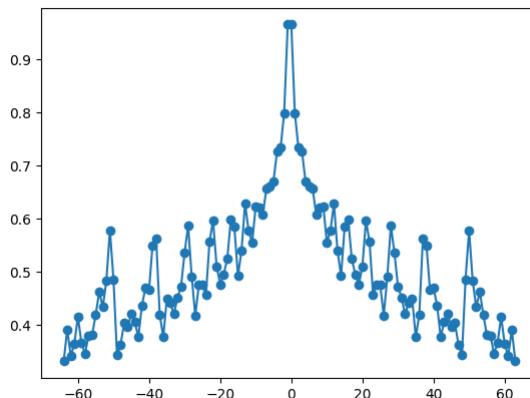
# Absolute Position Encoding

- Each position with a encoding vector
- E.g, sinusoidal

$$\vec{p}_i = [\sin \omega_1 i, \cos \omega_1 i, \dots, \sin \omega_{d/2} i, \cos \omega_{d/2} i]$$

- Why it helps

$$\vec{p}_i \cdot \vec{p}_j = \sum_{s=1}^{d/2} \cos \omega_s (i - j)$$



Imposes decaying for distant positions!



Neural Information Processing Systems  
<https://papers.neurips.cc/paper/7181-attention-is-all-you-need.pdf>

[Attention is All you Need](#)

by A Vaswani · Cited by 92585 — We propose a new simple network architecture, the Transformer, based solely on attention mechanisms, dispensing with recurrence and...  
11 pages

# Relative Position Encoding

- Some function of  $i - j$
- E.g., ALiBi's raw score

$$s_{i,j} = \frac{q_i \cdot k_j}{\sqrt{d}} - c|i - j|, c > 0$$



arXiv

<https://arxiv.org/pdf/> :

[Attention with Linear Biases Enables Input Length ...](#)

by O Press · 2021 · Cited by 150 – When using **ALiBi**, we do not add **positional embeddings** at the bottom of the network. ... CAPE: **encoding relative positions with continuous ...**

- Encourage to pay more attention to nearby tokens

# RoPE: Another Relative Position Encoding



arXiv

<https://arxiv.org> › pdf

⋮

## enhanced transformer with rotary position embedding

by J Su · 2021 · Cited by 303 — In this paper, we first investigate various methods to integrate positional information into the learning process of **transformer**-based language ...

- Rotate  $q_i$  by  $i \times \theta$ ,  $k_j$  by  $j \times \theta$
- So if  $|i - j|$  big, their inner product is small