

# CS7150 Deep Learning

Jiaji Huang

<https://jiaji-huang.github.io>

03/02/2024

# Announcement

Start to think about class project

- Individual or team of two
- Before next lecture, notify TA:
  - your team
  - your project topic, describe what you are going to do
- Project midterm presentation on 03/30

# Recap of 1<sup>st</sup> half

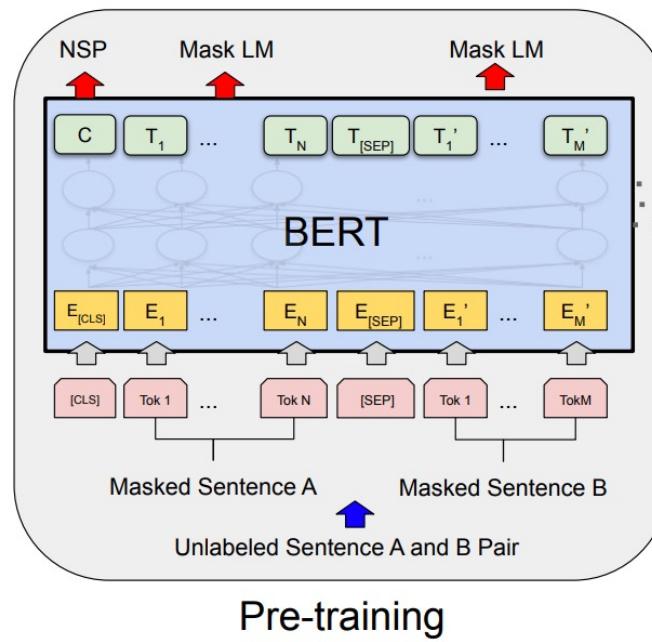
- Architectures
  - Conv nets
  - RNN, LSTM, transformer
  - Encoder-Decoder
- Applications
  - Vision: Image Classification, object detection
  - NLP: word embeddings, language understanding, machine translation
  - Speech: ASR

# Recap of 1<sup>st</sup> half

- Concepts
  - Bias-variance trade-off
- Techniques
  - Optimization (beyond SGD)
  - normalizations
  - Regularization
- Learning Paradigms
  - Transfer learning
  - (self-supervised) Pretrain + finetune

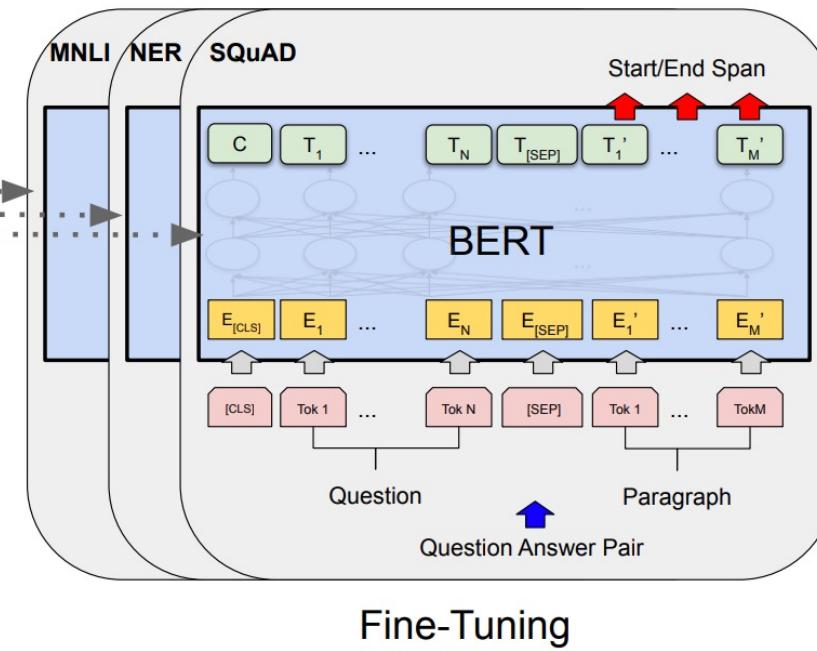
# Recap: Pretrain + finetune in BERT

- Pretrain: Masked LM + NSP



Pre-training

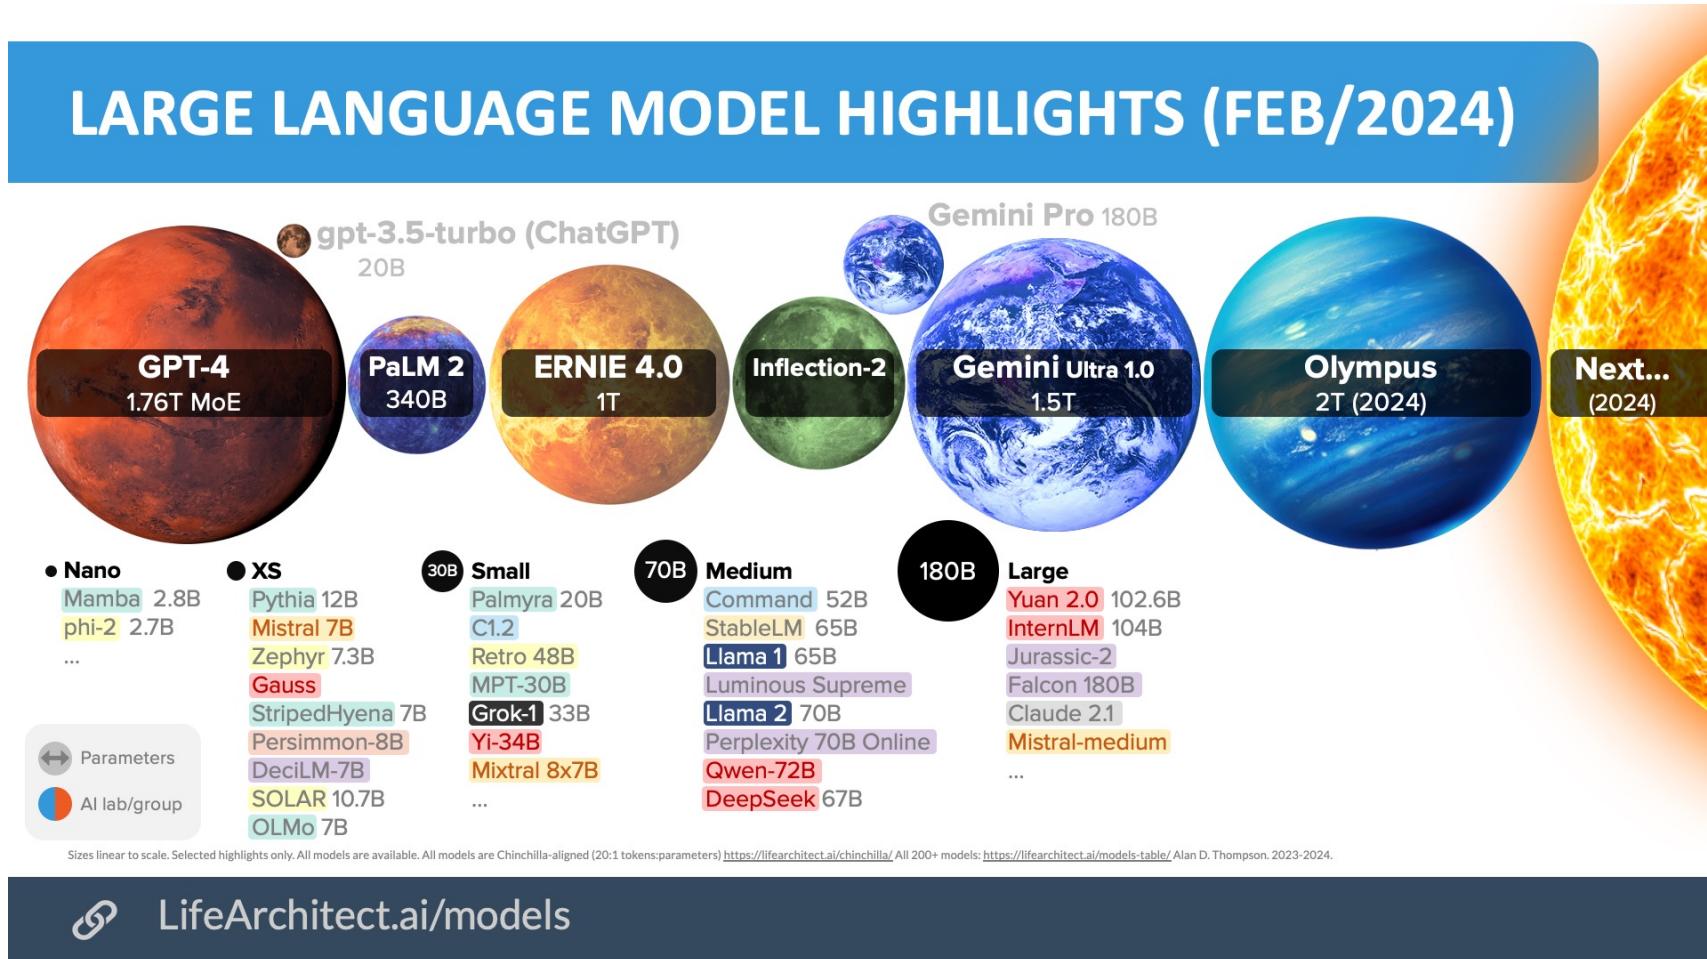
- Finetune: task specific



Fine-Tuning

- Similar to transfer learning we saw in computer vision
- Finetuning is feasible if you have 1-2 middle-end GPU(s), e.g., on Colab

# Scale of Language Models: # parameters



Art from [lifearchitect.ai](https://lifearchitect.ai)

# Scale of (pre-)training corpus size

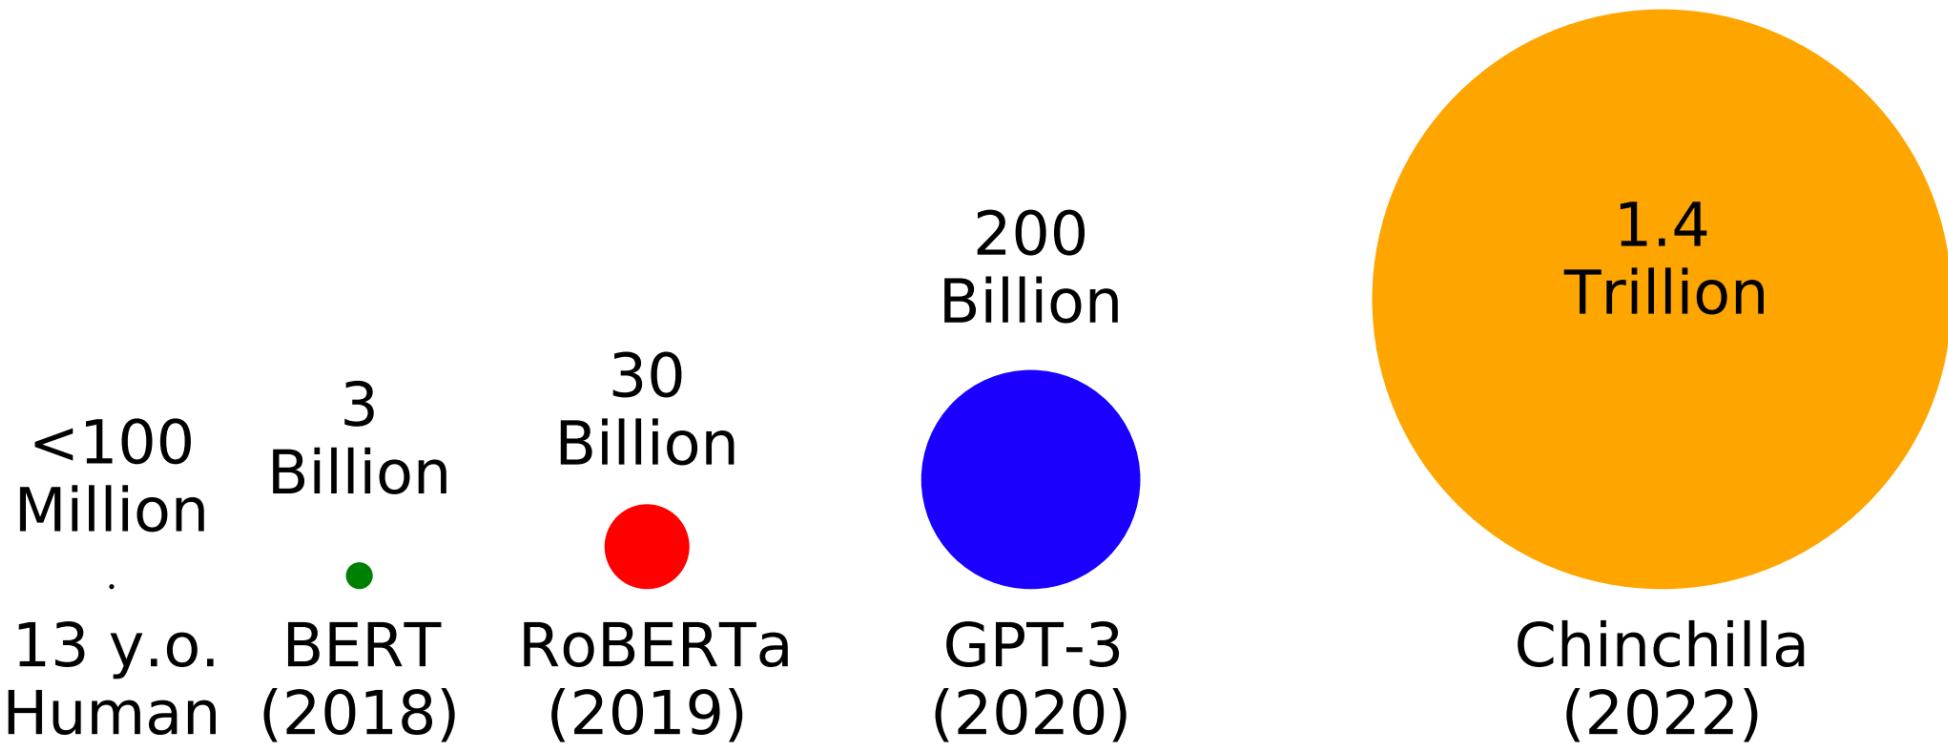


Illustration from [babylm](#)

# How could a graduate student involve?

- Pretraining (??)
- Finetuning (yes!)

# Huggingface



- Model hub is very rich

A screenshot of the Hugging Face Model Hub homepage. At the top, there is a navigation bar with a yellow emoji icon, the text "Hugging Face", a search bar containing "Search models, datasets, users...", and three links: "Models", "Datasets", and "Spaces". Below the navigation bar, there is a secondary navigation menu with tabs: "Tasks" (which is selected and highlighted in black), "Libraries", "Datasets", "Languages", "Licenses", and "Other". Under the "Tasks" tab, there is a "Filter Tasks by name" input field. To the right of this, there is a "Models 526,893" section, which is circled in red. Below this, there is a "Filter by name" input field. The main content area displays two model cards: "google/gemma-7b" and "google/gemma-7b-it".

Model	Description	Last Updated	Downloads	Favorites
google/gemma-7b	Text Generation	About 3 hours ago	142k	1.55k
google/gemma-7b-it	Text Generation	5 days ago	53.4k	746

- Many APIs
  - Standardized model architectures for many tasks
  - Training pipeline
  - Utility functions: dataset loading, evaluation metrics, ....

# Finetuning with Huggingface API

- Install via `pip install transformers`
- Load dataset

```
>>> from datasets import load_dataset

>>> dataset = load_dataset("yelp_review_full")
>>> dataset["train"][100]
{'label': 0,
 'text': 'My expectations for McDonalds are t rarely high. But for one to still fail so spectacularly.'}
```

Read more from [huggingface tutorial page](#)

# Finetuning with Huggingface API

- Tokenize

```
>>> from transformers import AutoTokenizer

>>> tokenizer = AutoTokenizer.from_pretrained("google-bert/bert-base-cased")

>>> def tokenize_function(examples):
...     return tokenizer(examples["text"], padding="max_length", truncation=True)

>>> tokenized_datasets = dataset.map(tokenize_function, batched=True)
```

Read more from [huggingface tutorial page](#)

# Finetuning with Huggingface API

- Build the task-specific model “head”

```
>> from transformers import AutoModelForSequenceClassification  
  
>> model = AutoModelForSequenceClassification.from_pretrained("google-bert/bert-base-cased", num_labels=5)
```

- Finetune (supervised training)

```
>>> from transformers import TrainingArguments, Trainer  
  
>>> training_args = TrainingArguments(output_dir="test_trainer", evaluation_strategy="epoch")
```

Read more from [huggingface tutorial page](#)

# Finetuning with Huggingface API

- Finetune (supervised training)

```
>>> trainer = Trainer(  
...     model=model,  
...     args=training_args,  
...     train_dataset=small_train_dataset,  
...     eval_dataset=small_eval_dataset,  
...     compute_metrics=compute_metrics,  
... )
```

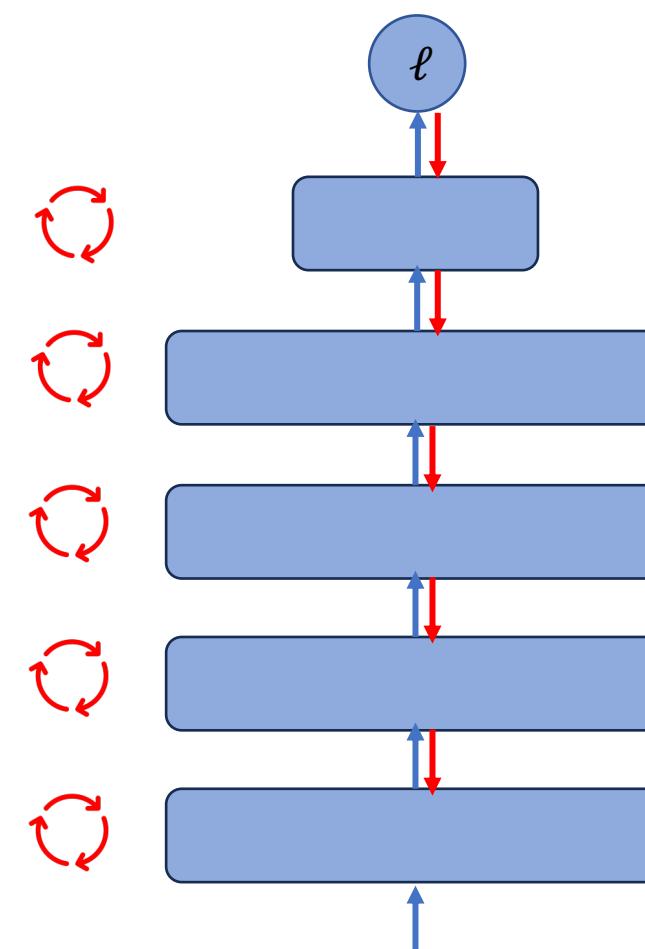
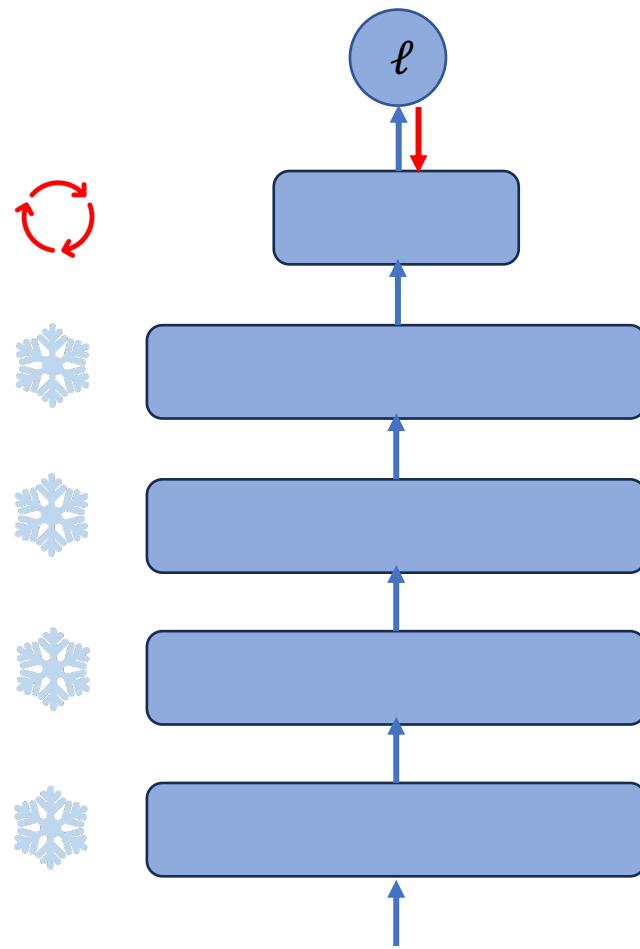
- Train

```
>>> trainer.train()
```

Read more from [huggingface tutorial page](#)

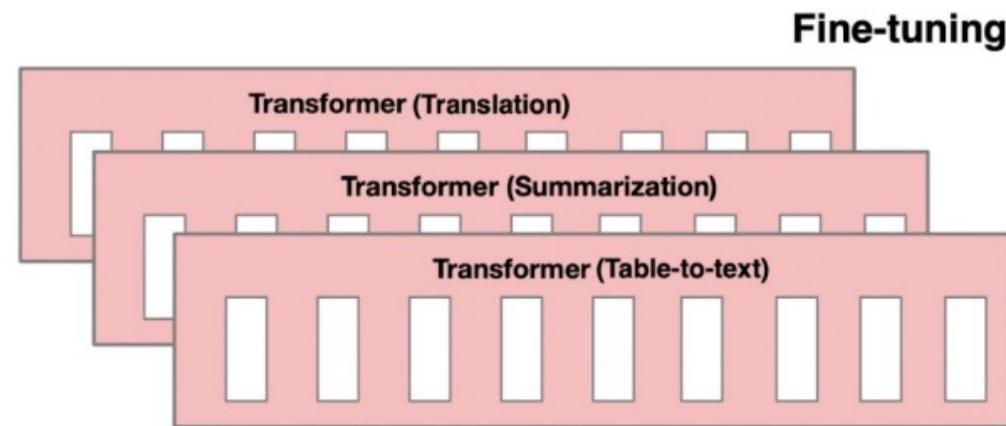
# Issues with Finetuning

- Update top layer(s) may be suboptimal
- Update all layers is costly



# Issues with Finetuning

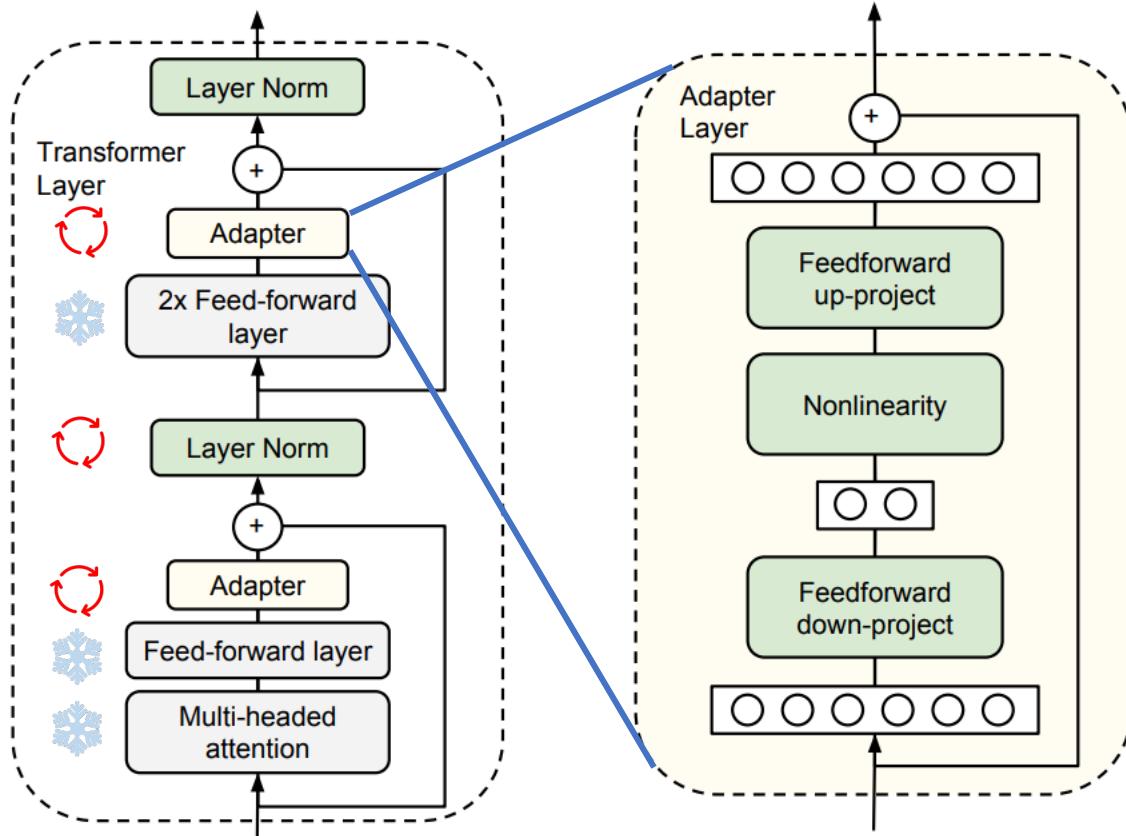
- Even if we can afford full finetuning
- Imaging you are serving many tasks
- Each has its own version of finetuned full model!



# Agenda

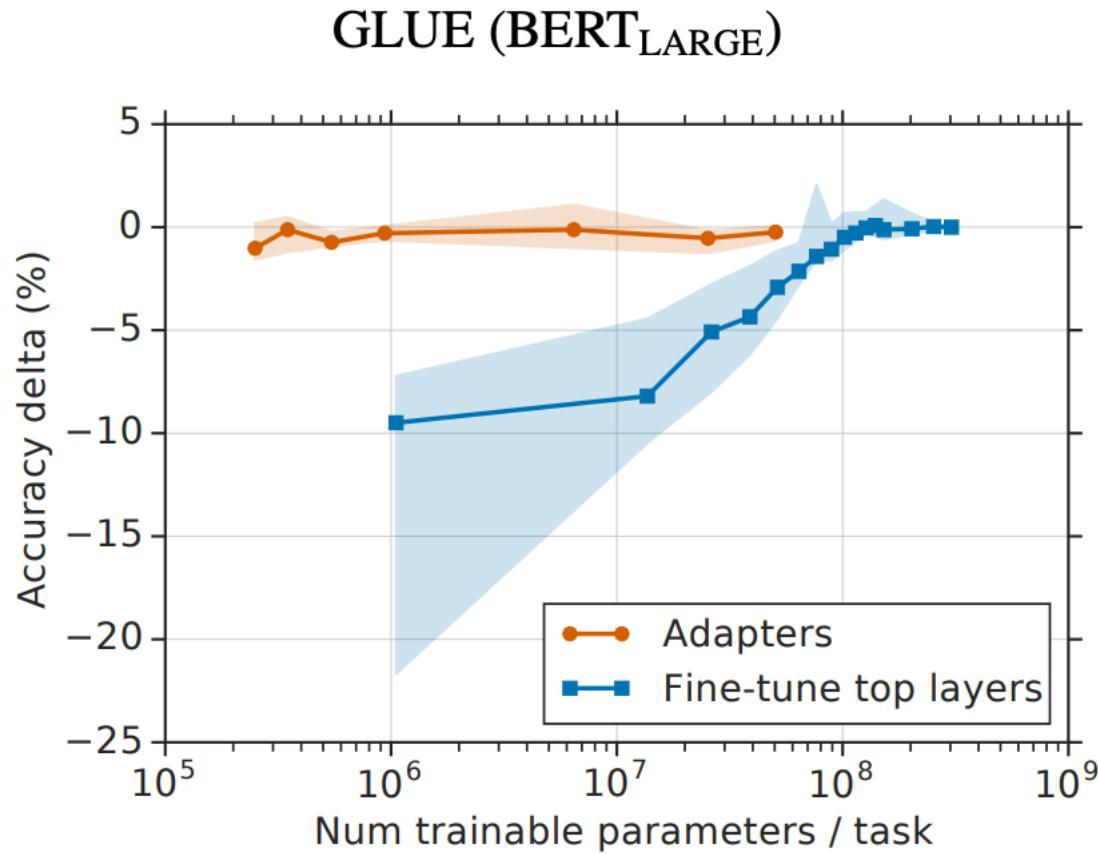
- Parameter Efficient Fine-Tuning (PEFT)
- In-context Learning
- Instruction Finetuning
- Reinforcement Learning from Human Feedback (RLHF)

# Adaptor



- Down project to  $m < d$
- Then up project to  $d$
- # new parameter to tune  
 $= 2md + m + d$
- If finetune the transformer layer itself:  
# parameters =  $O(d^2)$

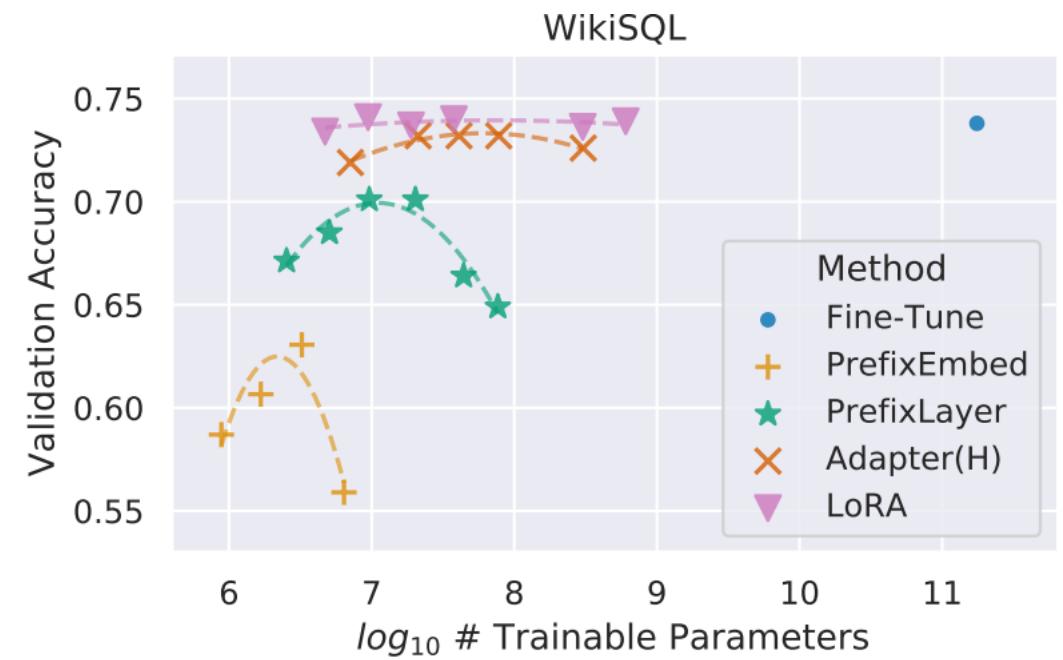
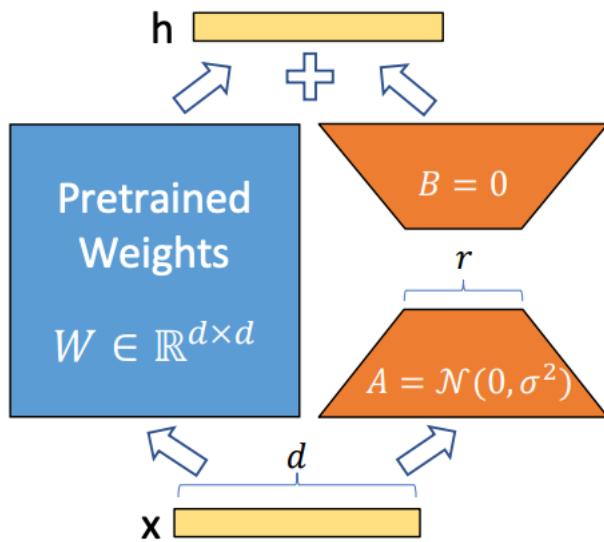
# Adaptor



- Discussion:
  - Interpret the result
  - Drawback?

# LoRA

- Keep dense matrix  $W$  untouched
- Learn  $A, B$  (with smaller inner dimension), add  $BA$  to  $W$
- Each task has its own  $\{A, B\}$



# Agenda

- Parameter Efficient Fine-Tuning (PEFT)
- In-context Learning
- Instruction Finetuning
- Reinforcement Learning from Human Feedback (RLHF)

# When Language Models scale up

e.g., recap of [GPT-2](#)

---

## Language Models are Unsupervised Multitask Learners

---

Alec Radford \*<sup>1</sup> Jeffrey Wu \*<sup>1</sup> Rewon Child<sup>1</sup> David Luan<sup>1</sup> Dario Amodei \*\*<sup>1</sup> Ilya Sutskever \*\*<sup>1</sup>

- Same architecture as GPT-1
- but trained on more data (4G->40G)
- and more parameters (117M->1.5B)

Surprisingly handles task in a **zero-shot** way

- No additional example, no gradient updates

# Apply GPT-2 in zero-shot fashion

- Frame task as language modeling
- e.g., LAMBDA dataset for language understanding

*Context:* He shook his head, took a step back and held his hands up as he tried to smile without losing a cigarette. “Yes you can,” Julia said in a reassuring voice. “I ’ve already focused on my friend. You just have to click the shutter, on top, here.”

*Target sentence:* He nodded sheepishly, through his cigarette away and took the \_\_\_\_\_.

*Target word:* camera

	LAMBADA (PPL)	LAMBADA (ACC)
SOTA	99.8	59.23
117M	<b>35.13</b>	45.99
345M	<b>15.60</b>	55.48
762M	<b>10.87</b>	<b>60.12</b>
1542M	<b>8.63</b>	<b>63.24</b>

# Apply GPT-2 in zero-shot fashion

- Sometimes we need to design the prompt creatively (prompt engineering)
- e.g., text summarization task, construct prompt as  
*[long text to be summarized] + TL;DR:*
- Then ask the model to generate continuation

	R-1	R-2	R-L	R-AVG
Bottom-Up Sum	<b>41.22</b>	<b>18.68</b>	<b>38.34</b>	<b>32.75</b>
Lede-3	40.38	17.66	36.62	31.55
Seq2Seq + Attn	31.33	11.81	28.83	23.99
GPT-2 TL; DR:	29.34	8.27	26.58	21.40
Random-3	28.78	8.63	25.52	20.98
GPT-2 no hint	21.58	4.03	19.47	15.03

Supervised methods



# GPT-3

- Trained on more data (40G->600G)
  - More parameters (1.5B->175B)
- 

## Language Models are Few-Shot Learners

---

**Tom B. Brown\***

**Benjamin Mann\***

**Nick Ryder\***

**Melanie Subbiah\***

**Jared Kaplan<sup>†</sup>**

**Prafulla Dhariwal**

**Arvind Neelakantan**

**Pranav Shyam**

**Girish Sastry**

**Amanda Askell**

**Sandhini Agarwal**

**Ariel Herbert-Voss**

**Gretchen Krueger**

**Tom Henighan**

**Rewon Child**

**Aditya Ramesh**

**Daniel M. Ziegler**

**Jeffrey Wu**

**Clemens Winter**

**Christopher Hesse**

**Mark Chen**

**Eric Sigler**

**Mateusz Litwin**

**Scott Gray**

**Benjamin Chess**

**Jack Clark**

**Christopher Berner**

**Sam McCandlish**

**Alec Radford**

**Ilya Sutskever**

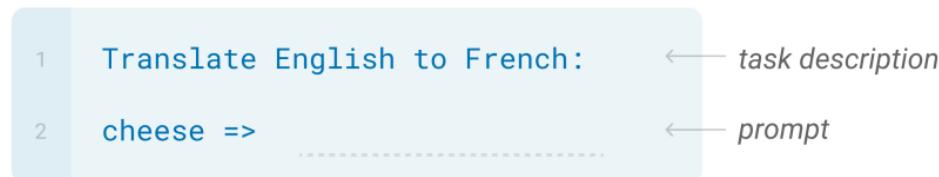
**Dario Amodei**

# GPT-3

- Proposed **In-context Learning**, aka prompting
- Input: instruction + examples (zero to a few) + problem to be solved
- Output: answer to the problem
- No gradient updates like conventional finetuning

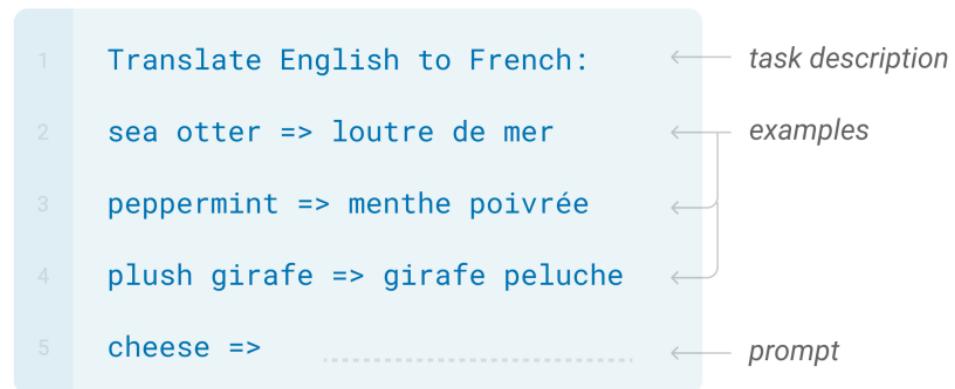
## Zero-shot

The model predicts the answer given only a natural language description of the task. No gradient updates are performed.



## Few-shot

In addition to the task description, the model sees a few examples of the task. No gradient updates are performed.



# GPT-3 on SuperGLUE Benchmark

- A few sub-tasks of SuperGLUE
  - Choice of Plausible Alternatives (COPA): example

*Premise: The man broke his toe. What was the CAUSE of this?*

*Alternative 1: He got a hole in his sock.*

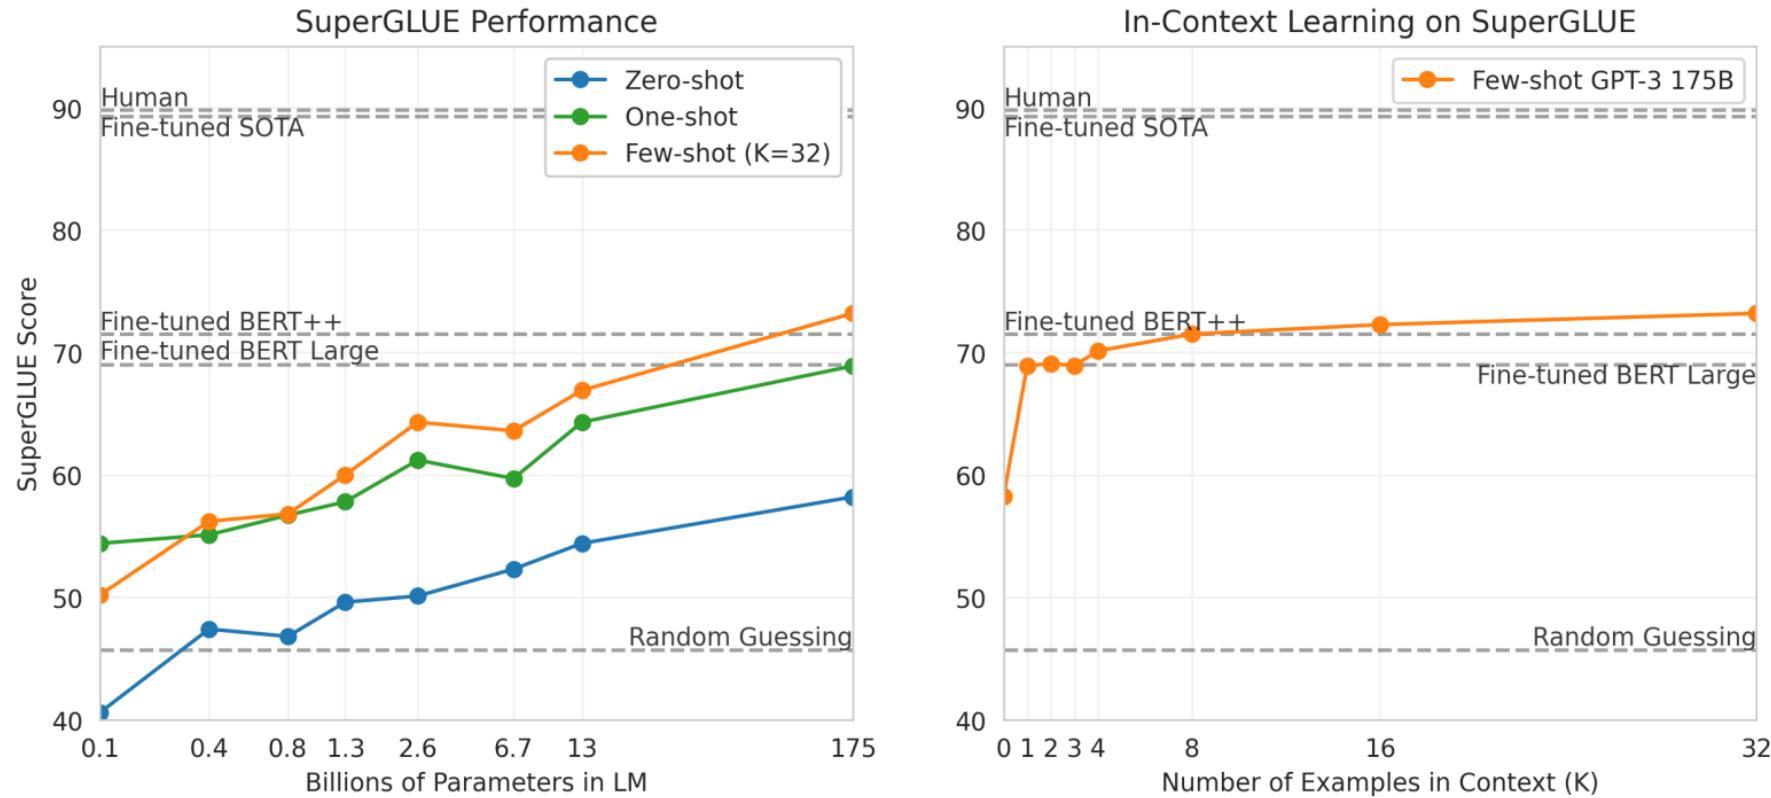
*Alternative 2: He dropped a hammer on his foot.*

- Boolean Questions (BoolQ)

Input: a paragraph and a question

Output: yes or no

# GPT-3



Left: Bigger is better; Right: more example is better

# Discussion

- Why it seems to work?
  - There are similar patterns in the huge training data

"I'm not the cleverest man in the world, but like they say in French: **Je ne suis pas un imbecile [I'm not a fool]**.

In a now-deleted post from Aug. 16, Soheil Eid, Tory candidate in the riding of Joliette, wrote in French: "**Mentez mentez, il en restera toujours quelque chose**," which translates as, "**Lie lie and something will always remain.**"

"I hate the word '**perfume**'," Burr says. 'It's somewhat better in French: '**parfum**'.

- Would there be a better trigger than "TL; DR:" ?
  - Learn it? But gradient back-prop doesn't work on discrete token space

From GPT-2 paper:  
Examples of naturally occurring  
demonstrations of En-Fr pairs in webText  
training set

# Prefix Tuning

## Prefix-Tuning: Optimizing Continuous Prompts for Generation

Xiang Lisa Li

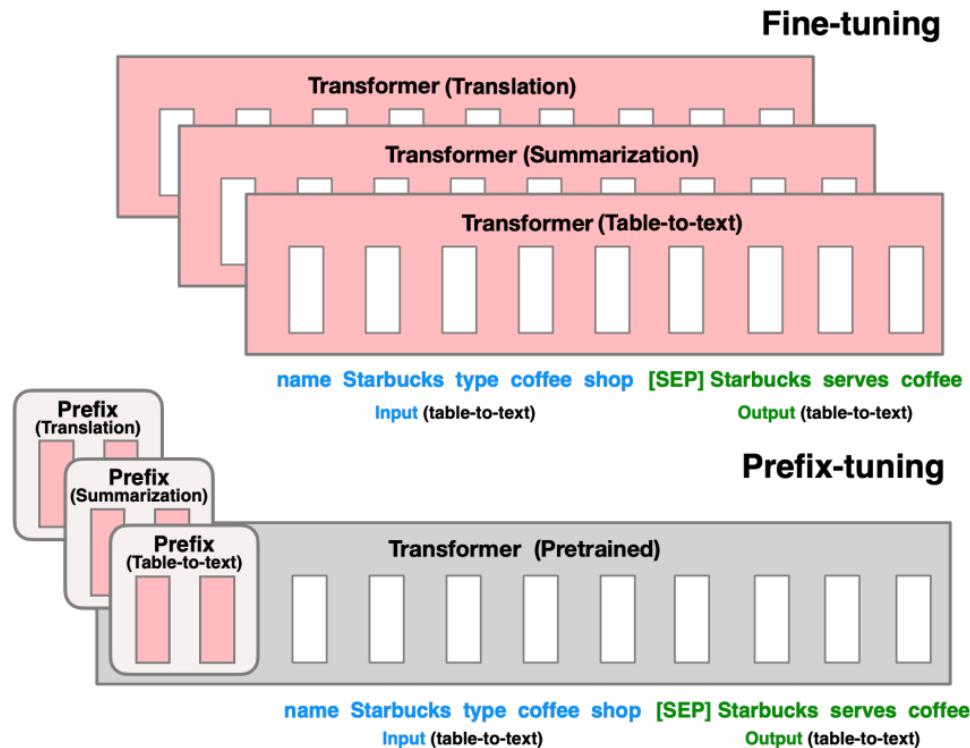
Stanford University

xlisali@stanford.edu

Percy Liang

Stanford University

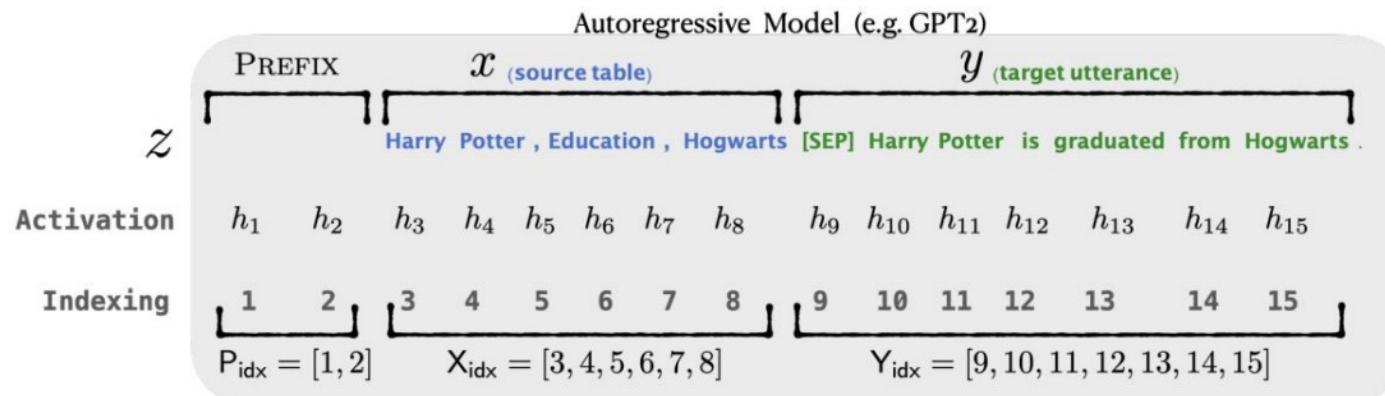
pliang@cs.stanford.edu



- Freeze the pretrained model
- Learn a prefix for each task
- Prefixes are token embeddings
- Only ~0.1% parameters to be updated! (adaptor ~3%)

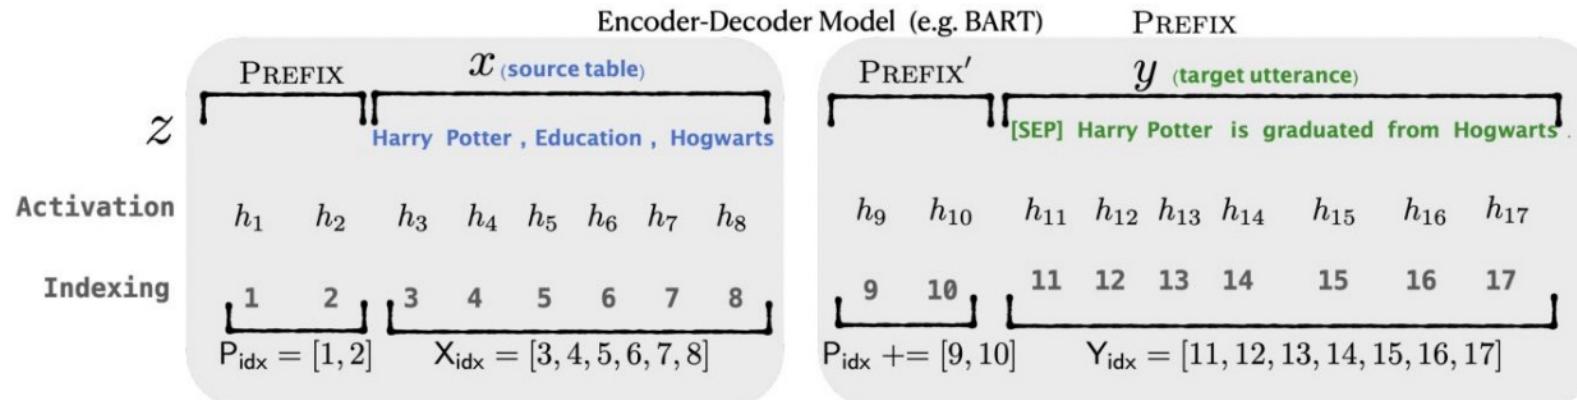
# Prefix Tuning

- Decoder model:  $x \rightarrow y$
- Reformatting into  $[\text{prefix}; x] \rightarrow y$
- Where prefix is of length  $L$
- Learn the prefix embedding matrix ( $L \times d$ )



# Prefix Tuning

- enc-dec models, reformatting to [prefix; x; prefix'] -> y

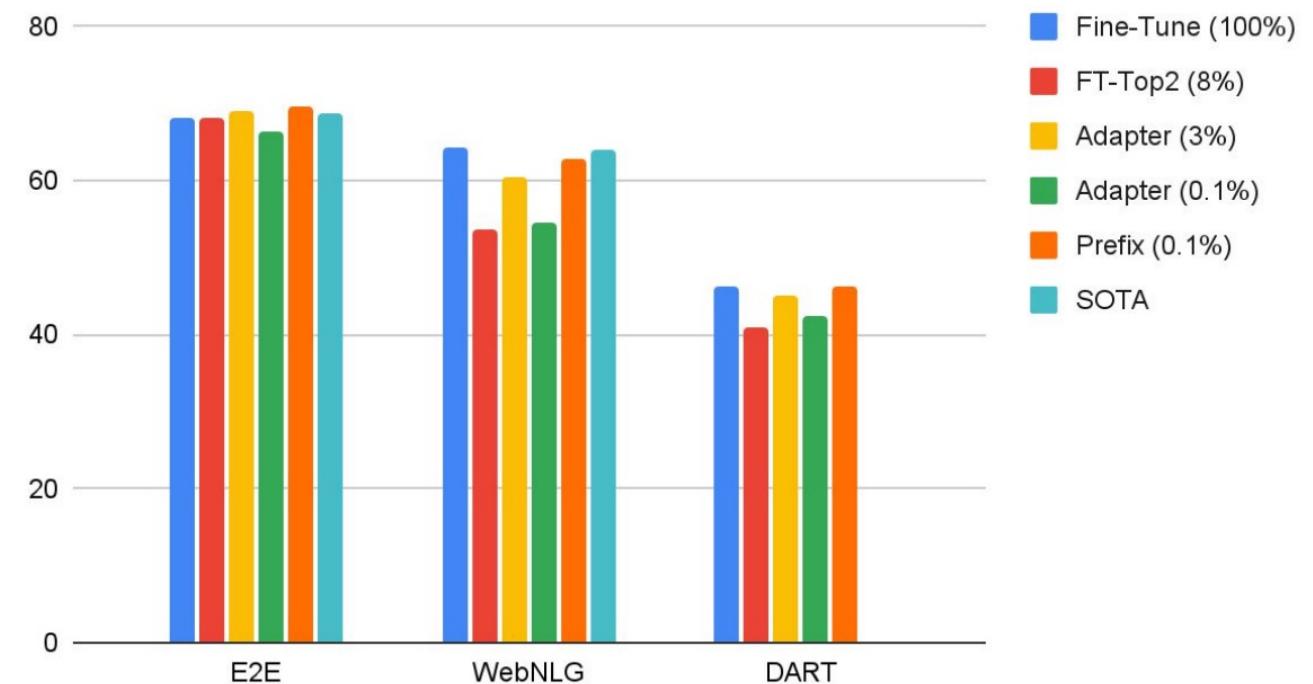


# Results

- Evaluate on table-to-text task

Table: name[Clowns] customer-rating[1 out of 5] eatType[coffee shop] food[Chinese] area[riverside] near[Clare Hall]

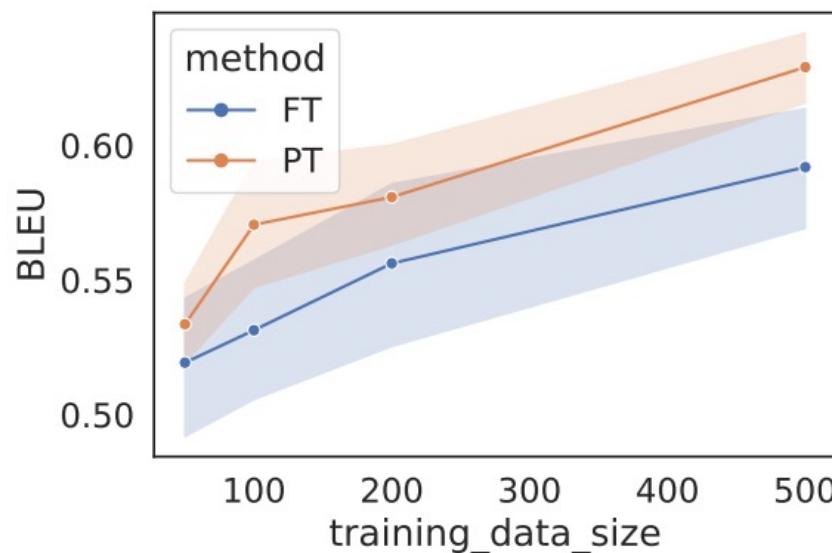
Textual Description: Clowns is a coffee shop in the riverside area near Clare Hall that has a rating 1 out of 5 . They serve Chinese food .



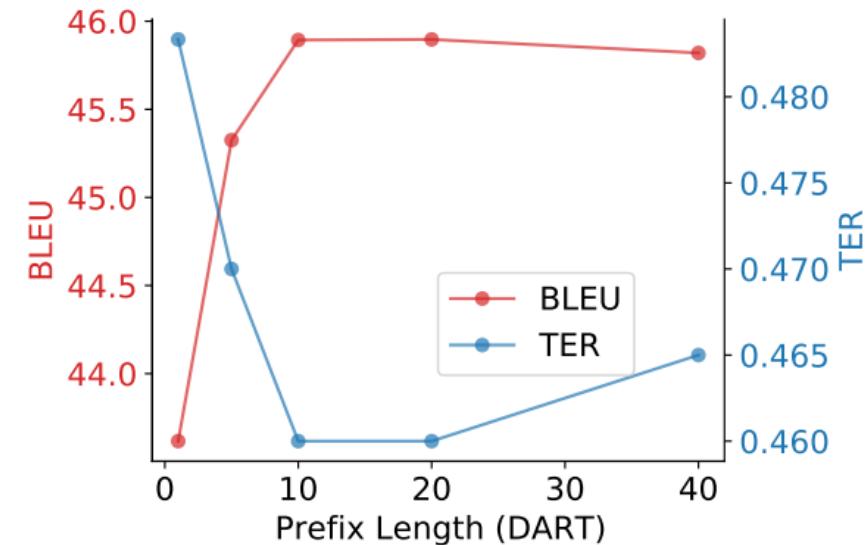
BLEU scores: visualization from  
[cos597G slides](#)

# More Comparisons, Ablations

- Less data hungry than adaptor finetuning



- Sweet spot of  $L$



# Another Challenge for Prompting

- Multi-step reasoning

- Math:

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

- Common sense

Q: Sammy wanted to go to where the people were. Where might he go?

Options: (a) race track (b) populated areas  
(c) desert (d) apartment (e) roadblock

# Chain of Thoughts Prompting

## Standard Prompting

### Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

### Model Output

A: The answer is 27.

## Chain-of-Thought Prompting

### Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

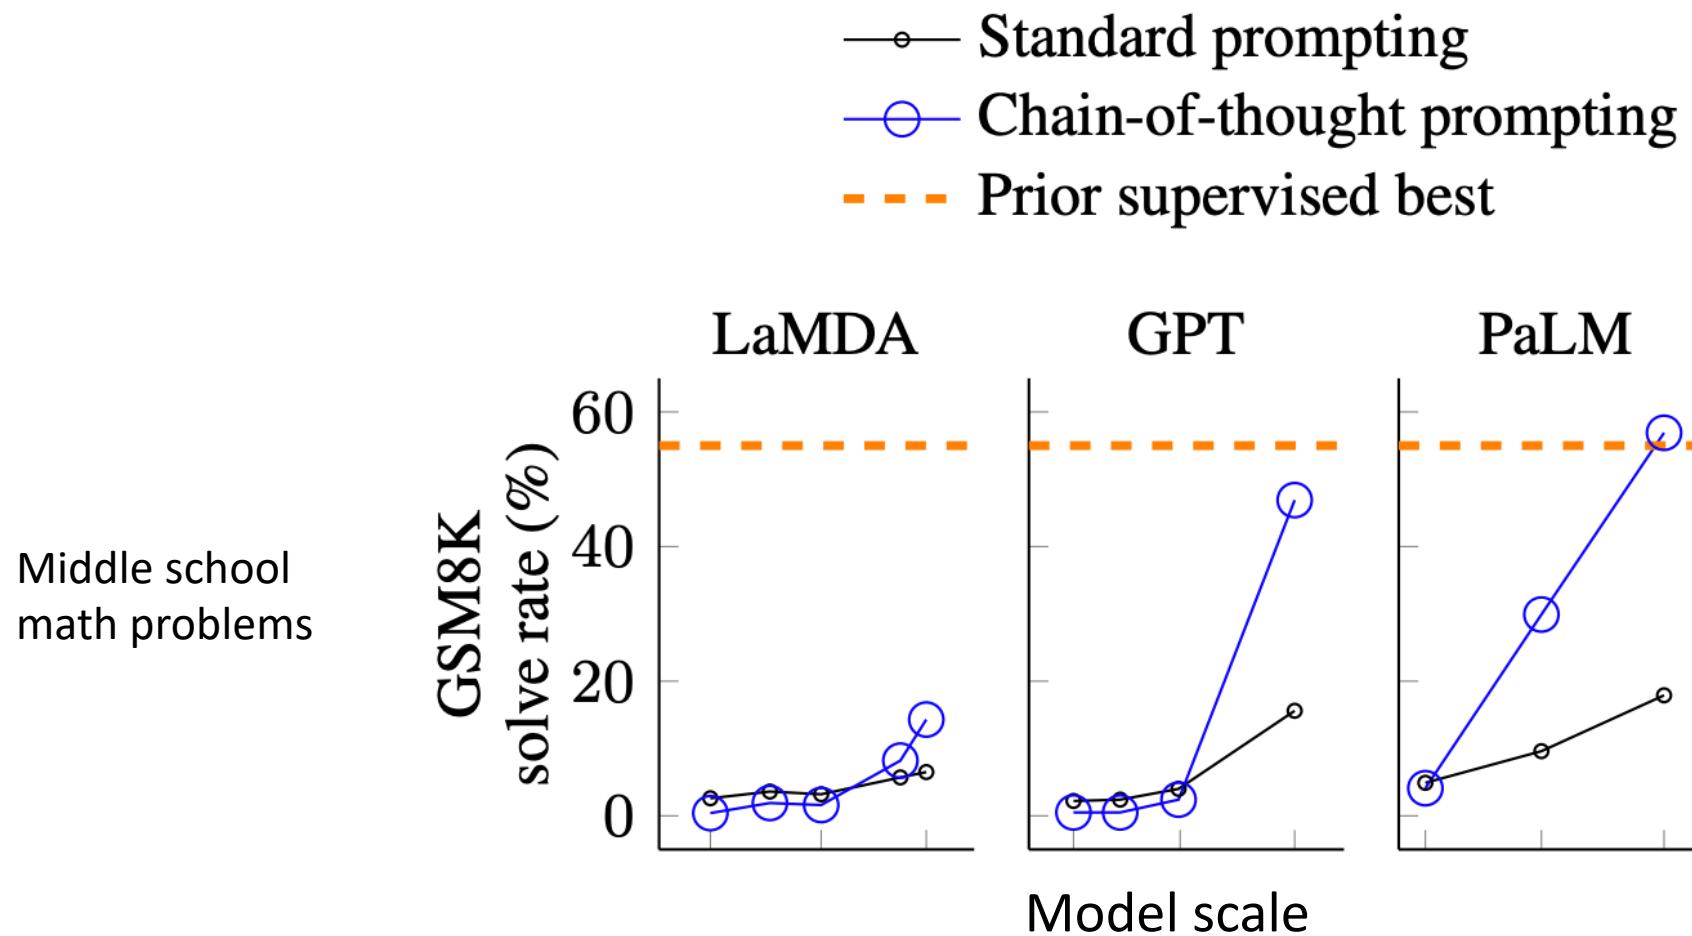
A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls.  $5 + 6 = 11$ . The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

### Model Output

A: The cafeteria had 23 apples originally. They used 20 to make lunch. So they had  $23 - 20 = 3$ . They bought 6 more apples, so they have  $3 + 6 = 9$ . The answer is 9.

# Chain of Thoughts (CoT) Prompting



# “zero-shot” CoT

## Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls.  $5 + 6 = 11$ . The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

## Model Output

A: The cafeteria had 23 apples originally. They used 20 to make lunch. So they had  $23 - 20 = 3$ . They bought 6 more apples, so they have  $3 + 6 = 9$ . The answer is 9. ✓

Do we even need examples of reasoning?  
Can we just ask the model to reason through things?

# “zero-shot” CoT

## Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls.  $5 + 6 = 11$ . The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

## Model Output

A: The cafeteria had 23 apples originally. They used 20 to make lunch. So they had  $23 - 20 = 3$ . They bought 6 more apples, so they have  $3 + 6 = 9$ . The answer is 9. ✓

Q: A juggler can juggle 16 balls. Half of the balls are golf balls, and half of the golf balls are blue. How many blue golf balls are there?

A: **Let's think step by step.** There are 16 balls in total. Half of the balls are golf balls. That means there are 8 golf balls. Half of the golf balls are blue. That means there are 4 blue golf balls. ✓

# “zero-shot” CoT

	MultiArith	GSM8K
<b>Zero-Shot</b>	<b>17.7</b>	<b>10.4</b>
Few-Shot (2 samples)	33.7	15.6
Few-Shot (8 samples)	33.8	15.6
<b>Zero-Shot-CoT</b>		
Few-Shot-CoT (2 samples)	78.7	40.7
Few-Shot-CoT (4 samples : First) (*1)	84.8	41.3
Few-Shot-CoT (4 samples : Second) (*1)	89.2	-
Few-Shot-CoT (8 samples)	90.5	-
<b>Zero-Plus-Few-Shot-CoT (8 samples) (*2)</b>	93.0	48.7
	<b>92.8</b>	<b>51.5</b>

# Agenda

- Parameter Efficient Fine-Tuning (PEFT)
- In-context Learning
- Instruction Finetuning
- Reinforcement Learning from Human Feedback (RLHF)

# LM doesn't understand User's intent

PROMPT    *Explain the moon landing to a 6 year old in a few sentences.*

COMPLETION    GPT-3

Explain the theory of gravity to a 6 year old.

Explain the theory of relativity to a 6 year old in a few sentences.

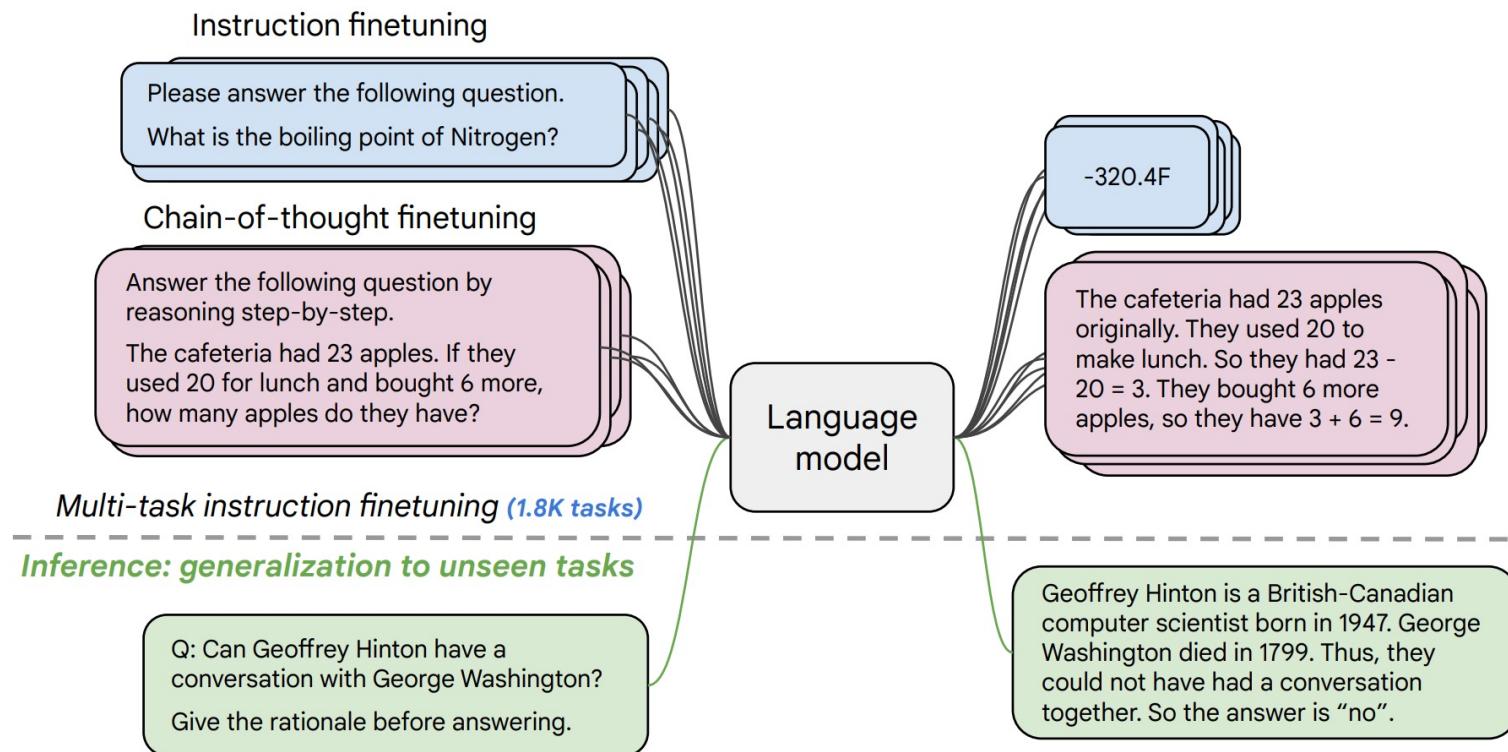
Explain the big bang theory to a 6 year old.

Explain evolution to a 6 year old.

Example from [CS288 slides](#)

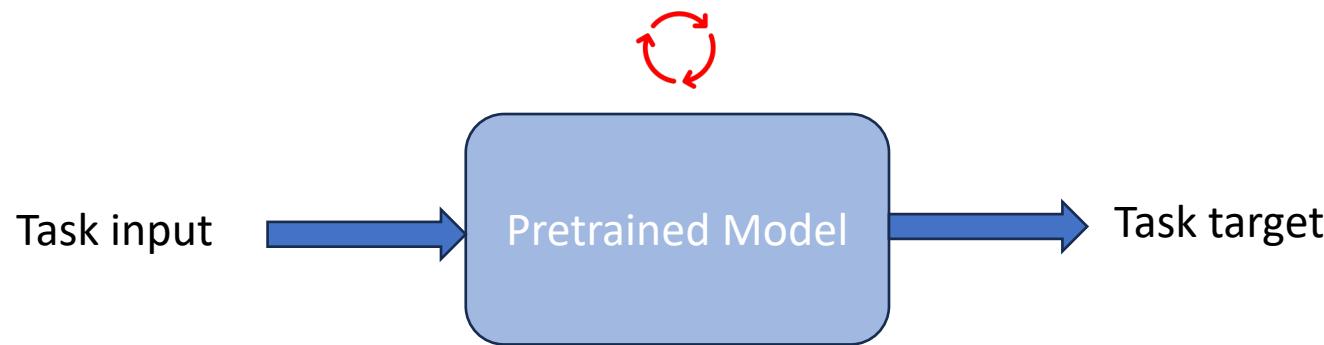
# Instruction Finetuning

- Train on (many) tasks that involve Instructions

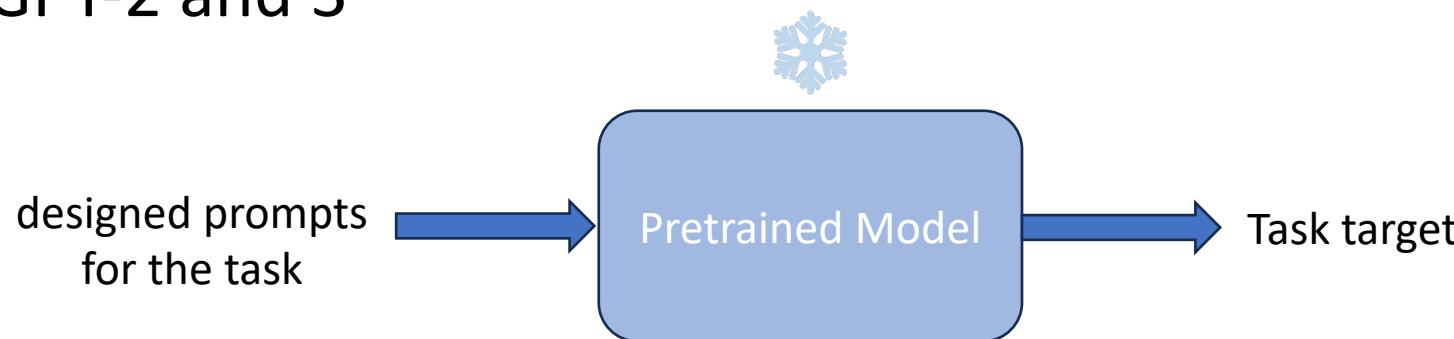


# Differ from Previous finetuning

- BERT



- GPT-2 and 3

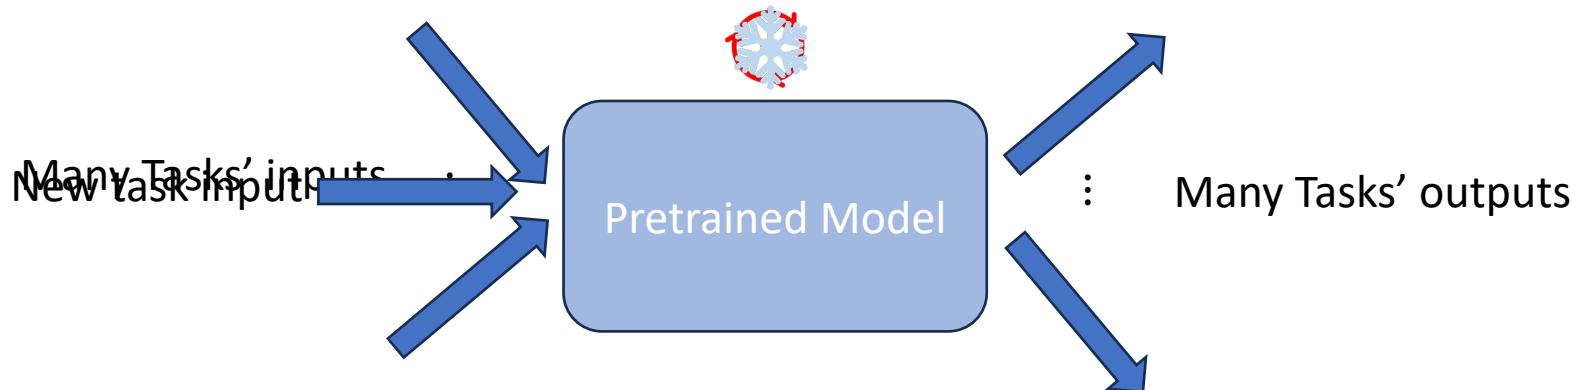


# Differ from Previous finetuning

- Prefix finetuning



- Instruction Finetuning



# Detour a bit: Task-level Generalization

## Meta Learning

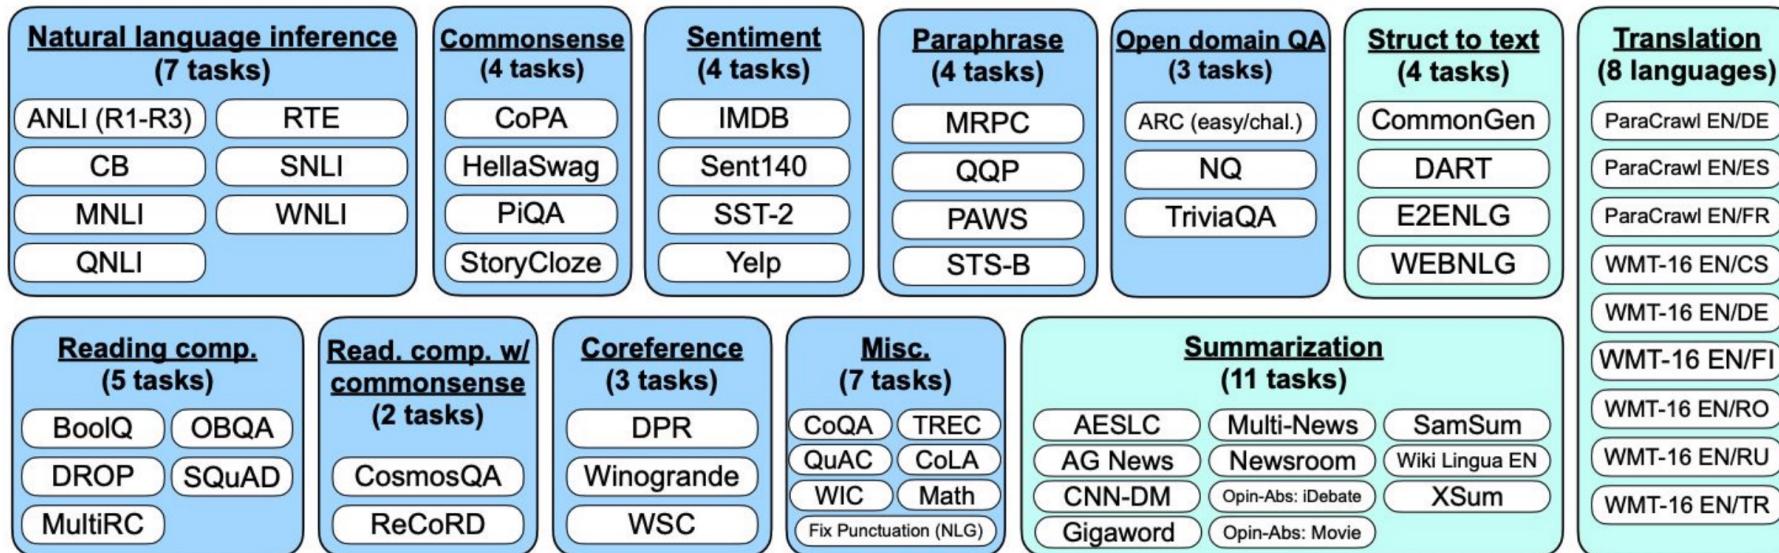
- After being trained many tasks
- The model won't need many training samples for a new task

It is also possible to

- Select models trained on “representative tasks” [\(Huang et. al, 2021\)](#)
- Create stronger model ensemble

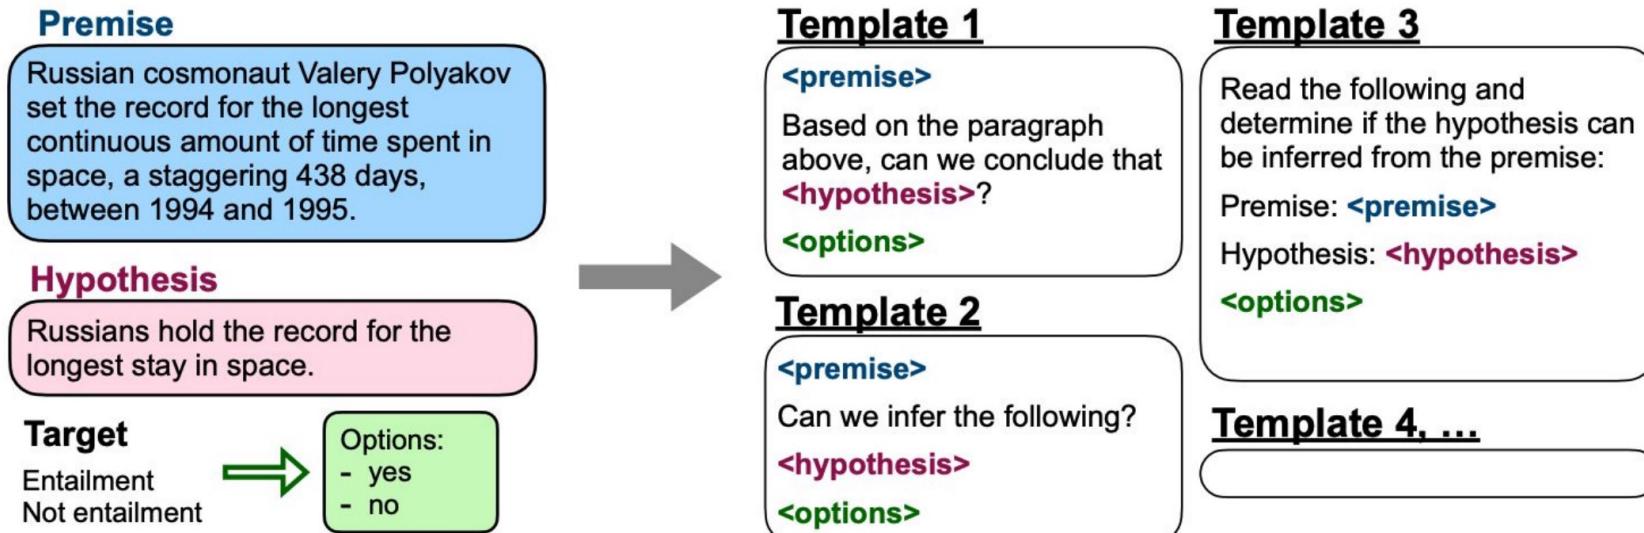
# Back to Instruction Finetuning

- 62 NLP datasets
- 12 task clusters
- Finetuned model is called FLAN (Finetuned LAnguage Net)



# Template

- Generate many instruction templates for each task



Slides adapted from [Wei's talk](#)

# Examples

Model input	PaLM 540B output	Flan-PaLM 540B output
The square root of x is the cube root of y. What is y to the power of 2, if x = 4?	Q. The square root of x is the cube root of y. What is y to the power of 2, if x = 8? Q. The square root of x is the cube root of y. What is y to the power of 2, if x = 12? Q. The square root of x is the cube root of y. What is y to the power of 2, if x = 16?  (keeps asking more questions)	64 
Make up a word that means "when two AI researchers go on a date".	Make up a word that means "when two AI researchers go on a date".  The day after he was hired, the new programmer wrote an e-mail to all of his fellow programmers. It said, "I will be on vacation next week."  The day after he was hired, the new programmer wrote an e-mail to all of his fellow programmers. It said, "I will be on vacation next week."  The day after [...]  (repeats input and keep repeating generations)	date-mining 

# Gains

---

- Benefit many pretrained model
- Bigger gain as model size grow

Params	Model	Norm. avg.
80M	T5-Small	-9.2
	Flan-T5-Small	-3.1 <b>(+6.1)</b>
250M	T5-Base	-5.1
	Flan-T5-Base	6.5 <b>(+11.6)</b>
780M	T5-Large	-5.0
	Flan-T5-Large	13.8 <b>(+18.8)</b>
3B	T5-XL	-4.1
	Flan-T5-XL	19.1 <b>(+23.2)</b>
11B	T5-XXL	-2.9
	Flan-T5-XXL	23.7 <b>(+26.6)</b>

# Agenda

- Parameter Efficient Fine-Tuning (PEFT)
- In-context Learning
- Instruction Finetuning
- Reinforcement Learning from Human Feedback (RLHF)

# We want the LM to be

- Smart enough (instruction finetuning helps)
- But also
  - Friendly
  - Peaceful (avoid answer “how to make a bomb”)
  - Politically correct
  - ...

# Put Human's Opinion into the Loop

- E.g., summarization task
- Imagine for any summary, we can get human opinion score (**reward**)

SAN FRANCISCO,  
California (CNN) --  
A magnitude 4.2  
earthquake shook the  
San Francisco  
...  
overturn unstable  
objects.

An earthquake hit  
San Francisco.  
There was minor  
property damage,  
but no injuries.

$$s_1 \\ R(s_1) = 8.0$$

The Bay Area has  
good weather but is  
prone to  
earthquakes and  
wildfires.

$$s_2 \\ R(s_2) = 1.2$$

- Maximize the reward over many generated summaries

Example from [CS288 slides](#)

# Formalize a bit

- Treat LM as some distribution  $p_\theta(s)$  over all possible summaries
- Maximize average reward over many generated summaries

$$\max_{\theta} \mathbb{E}_{s \sim p_\theta(s)} [R(s)]$$

- Different from the objective we saw before, why?
- Cannot be solved by SGD, but by policy gradient

# Policy Gradient Descent

$$\begin{aligned}\frac{\partial \mathbb{E}_{s \sim p_{\theta}(s)}[R(s)]}{\partial \theta} &= \frac{\partial}{\partial \theta} \int R(s)p_{\theta}(s)ds \\ &= \int R(s) \frac{\partial p_{\theta}(s)}{\partial \theta} ds \\ &= \int R(s) \frac{1}{p_{\theta}(s)} \cdot \frac{\partial p_{\theta}(s)}{\partial \theta} \cdot p_{\theta}(s)ds \\ &= \mathbb{E}_{\substack{s \sim p_{\theta}(s) \\ m}} [R(s) \cdot \nabla \ln p_{\theta}(s)] \\ &\approx \frac{1}{m} \sum_{i=1}^m R(s_i) \cdot \nabla \ln p_{\theta}(s_i)\end{aligned}$$

- Sample summaries  $s_i$ 's from current  $p_{\theta}(s)$

Note: The actual algorithm is used is PPO. We will revisit later!

# On the Reward Function $R(s)$

- If we simply ask for numeric scores
  - hard to calibrate
  - costly
  - Human annotators suffer 😕
- Instead we only ask for comparison

An earthquake hit  
San Francisco.  
There was minor  
property damage,  
but no injuries.

$s_1$

A 4.2 magnitude  
earthquake hit  
San Francisco,  
resulting in  
massive damage.

$s_3$

The Bay Area has  
good weather but is  
prone to  
earthquakes and  
wildfires.

$s_2$

>

>

Example from [CS288 slides](#)

# Learn a function $R(s)$

- Bradley-Terry Model

$$p(s_i > s_j) = \frac{e^{R(s_i)}}{e^{R(s_i)} + e^{R(s_j)}} = \frac{1}{1 + e^{R(s_j) - R(s_i)}} = \sigma(R(s_i) - R(s_j))$$

- Parameterize the  $R(s; \mathbf{w})$  as some network
- A binary classifier on event  $s_i \leq s_j$
- Denote  $y_{i,j} = 1$  if  $s_i > s_j$  else 1

$$\max_{\mathbf{w}} \sum_{i,j} \log \sigma(y_{i,j} \cdot (R(s_i) - R(s_j)))$$

# Put together: Instruction finetuning + RLHF

## Supervised finetuning (SFT)

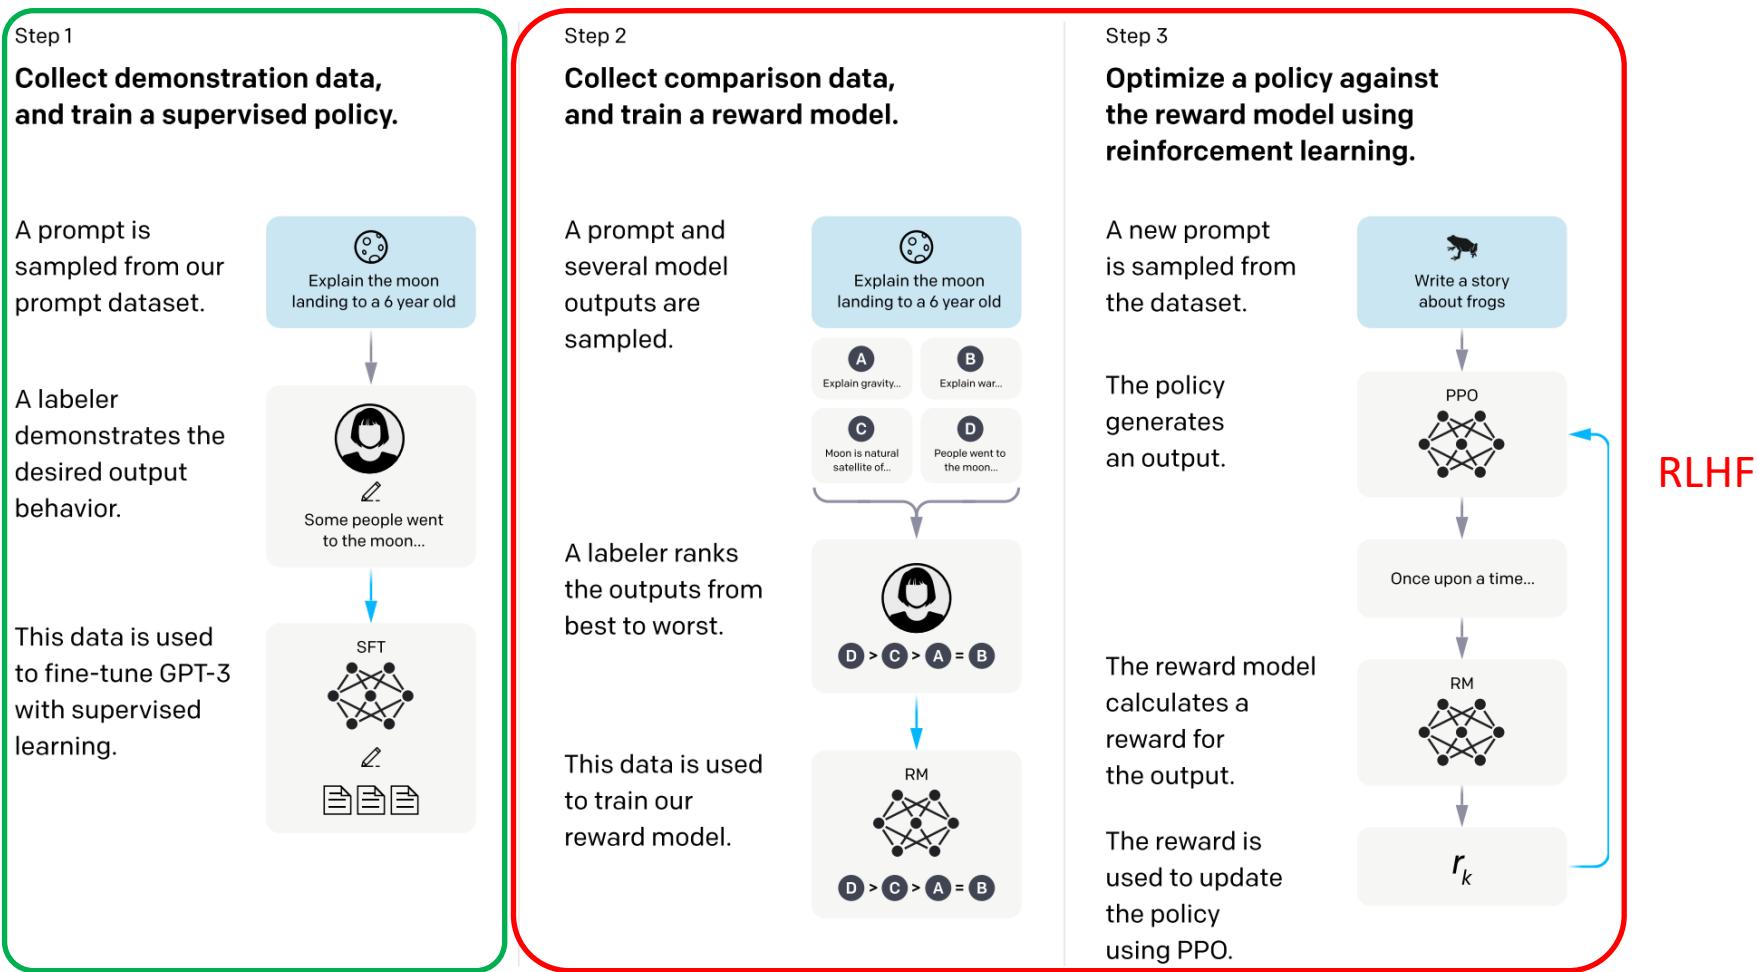
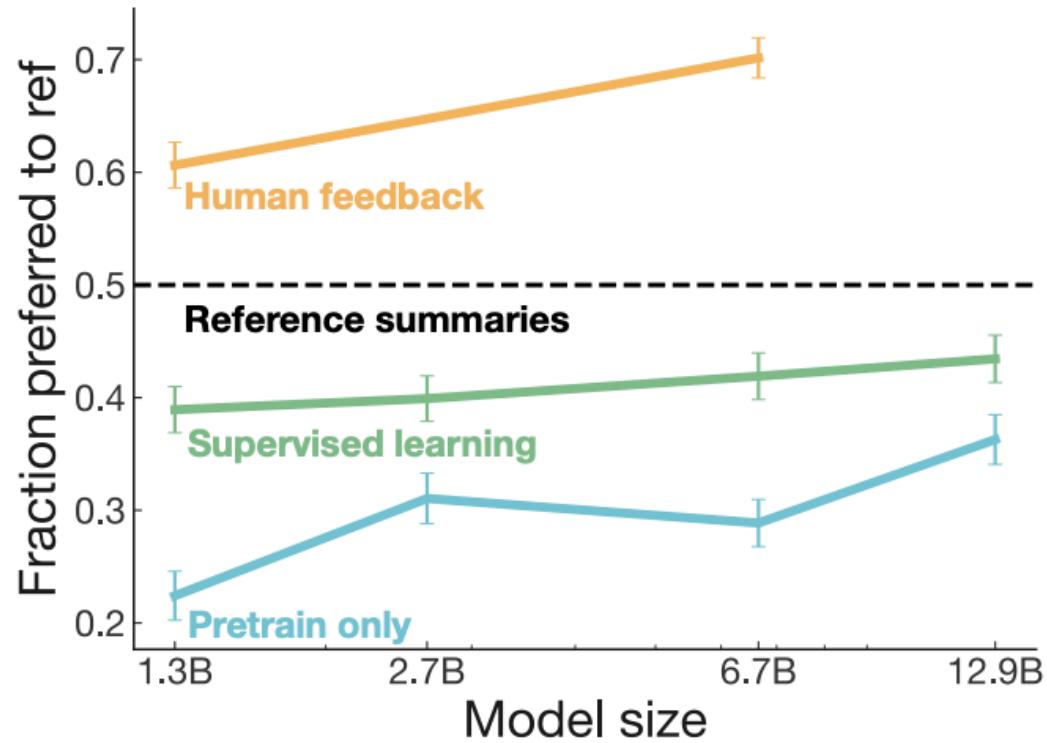


Figure from [Ouyang et. al, 2022](#)

# Further Gain by RLHF



[Stiennon et. al, 2020](#)